

BIG DATA EUROPE

Empowering Communities  
with Data Technologies

# BIG DATA EUROPE

## Introduction: Project, Architecture, Components, and Interfaces



Ivan Ermilov

University of Leipzig/InfAI

[iermilov@informatik.uni-leipzig.de](mailto:iermilov@informatik.uni-leipzig.de)

09 October 2017

ICTCS, Amman, Jordan



# Talk outline

2

- ⊙ The BigDataEurope action
- ⊙ The Big Data Integrator platform
  - Technical Architecture
  - Components & Interfaces
- ⊙ Pilots across all seven H2020 challenges

# BigDataEurope Action & Platform



# Big Data Europe (CSA: 2015-17)

- ⊙ **Show** societal value of Big Data
  - Across all societal challenges addressed by Horizon 2020
- ⊙ **Lower** barrier for using big data technologies
  - Effort and resources to convert tools and workflows
  - Skills and expertise
- ⊙ **Help** establish data value chains
  - Across languages, organizations, and domains



# BigDataEurope: Summary

5



Health



Food & Agriculture



Energy



Transport



Climate



Social Sciences



Security

- ⊙ Horizon2020 project
- ⊙ 17 partners
- ⊙ 7 pilots in various domains
- ⊙ > 30 Big Data components
- ⊙ > 250 stars on github



# Consortium

6



NCSR  
DEMOKRITOS




Food and Agriculture  
Organization of the  
United Nations



- ⊙ Semantic Web and Linked Data expertise
- ⊙ SW/LD technologies



# Stakeholder Engagement



## Stakeholder engagement workshops:

- ⊙ Present the action and its showcase deployments
- ⊙ Raise awareness about BDE results and what they mean for stakeholders
- ⊙ Collect requirements to drive further development



# Big Data Integrator



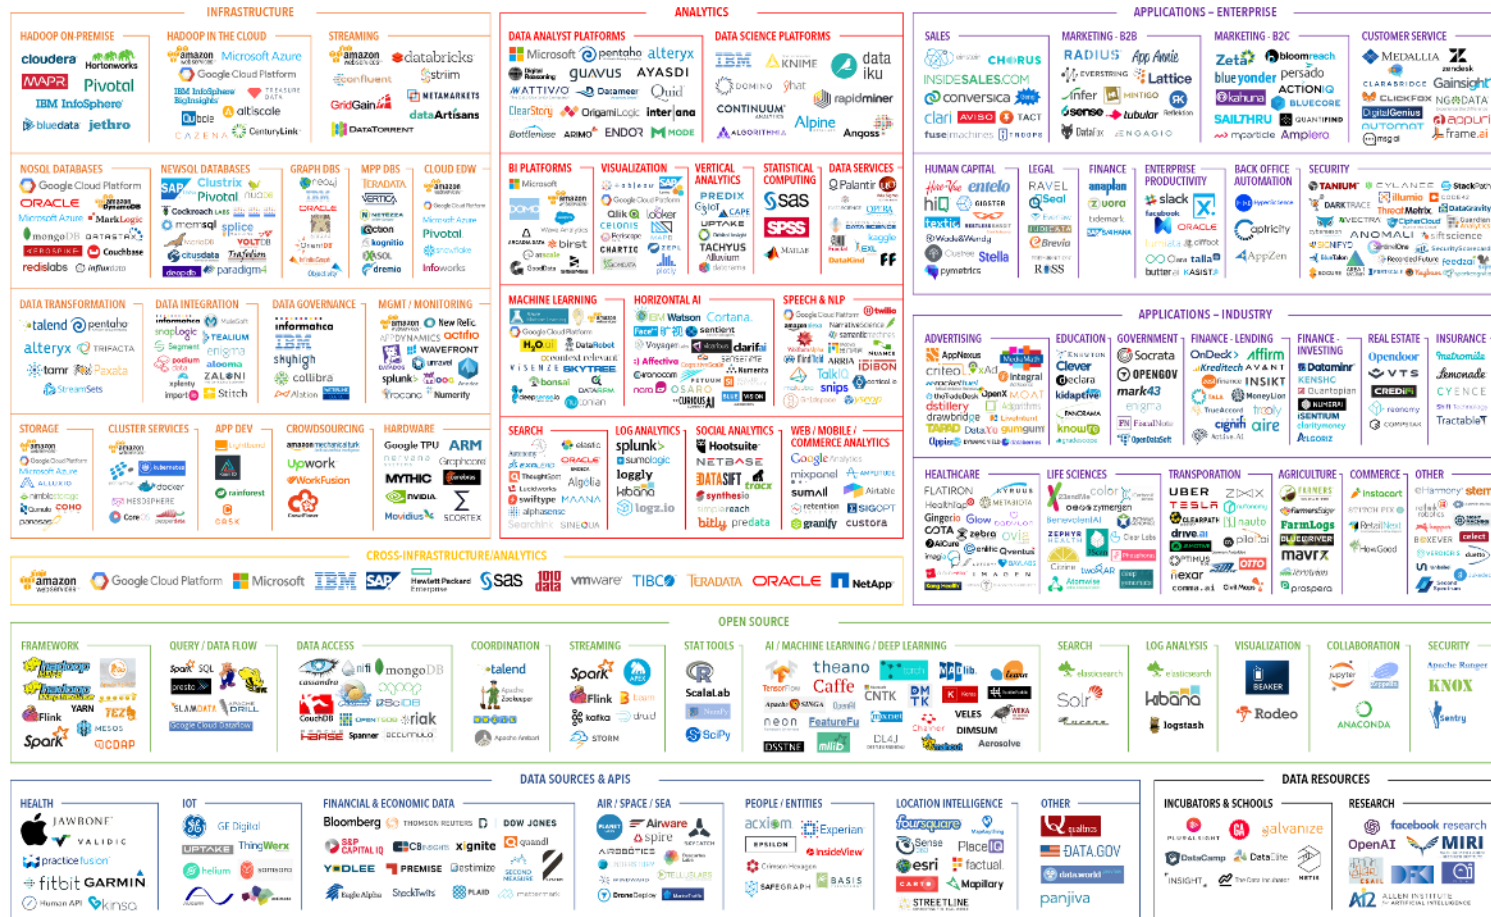


# Platform Goals

9

- ⊙ Open source
- ⊙ Simple to get started with Big Data
- ⊙ Support a variety of use cases
- ⊙ Embrace emerging Big Data technologies
- ⊙ Simple integration with custom components

## BIG DATA LANDSCAPE 2017





# Architecture

11

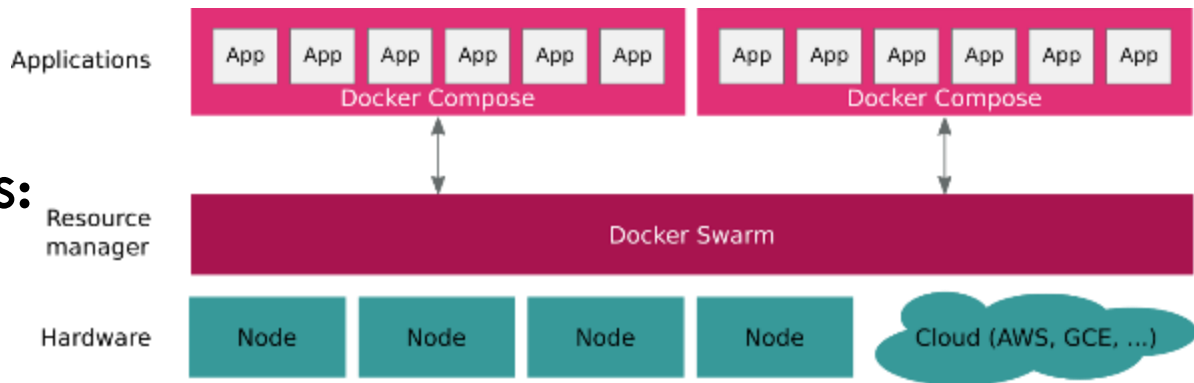
- ◎ Big Data Integrator (BDI):
  - The prototype developed by BDE
- ◎ Main points of the architecture
  - Dockerization
  - Support layer, including integrated UI
  - Semantification layer



# Docker containers

12

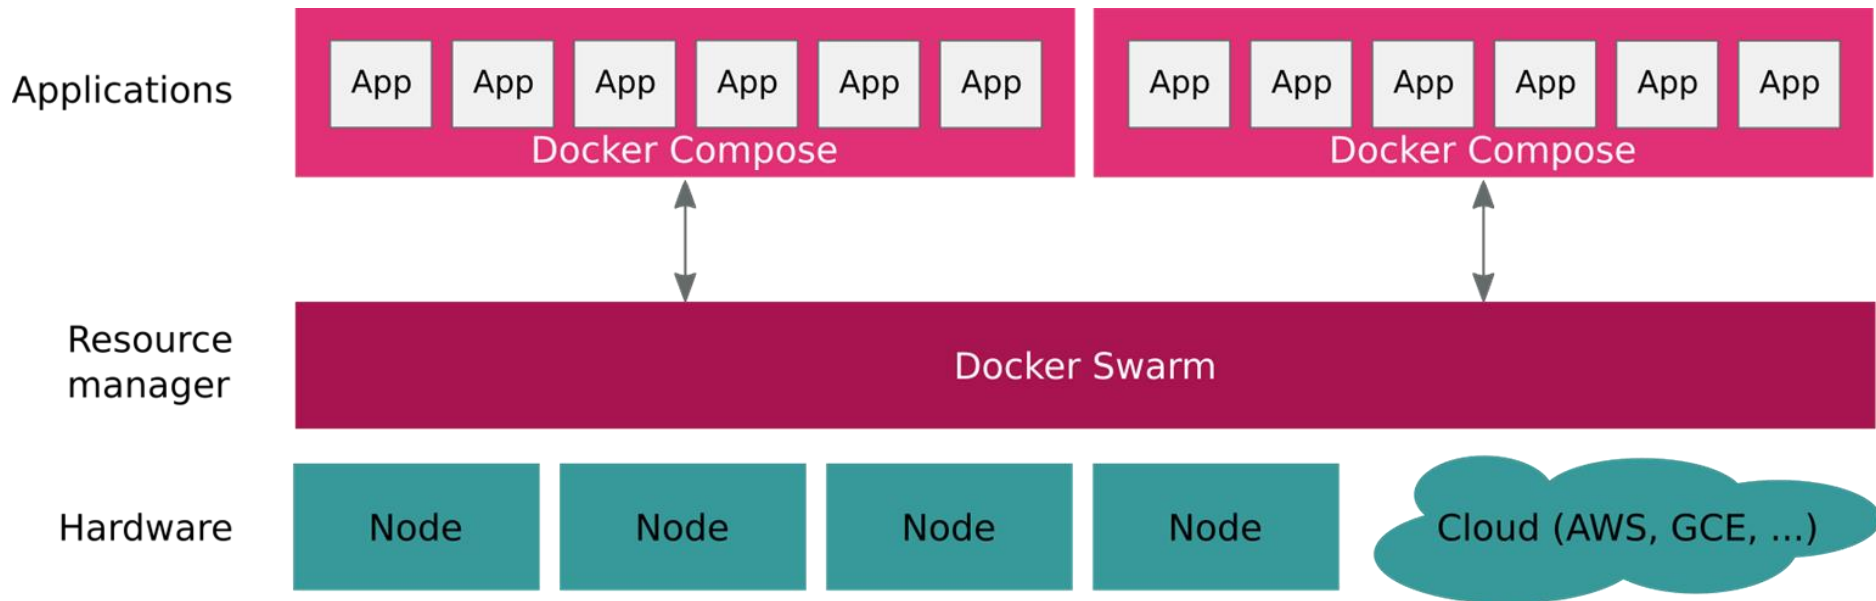
- ⊙ Docker offers lightweight virtualization
  - Docker containers can be shared to be provisioned on different Linux variations and versions
- ⊙ Identical base sys not required
- ⊙ All BDI components: Docker containers





# Platform Architecture

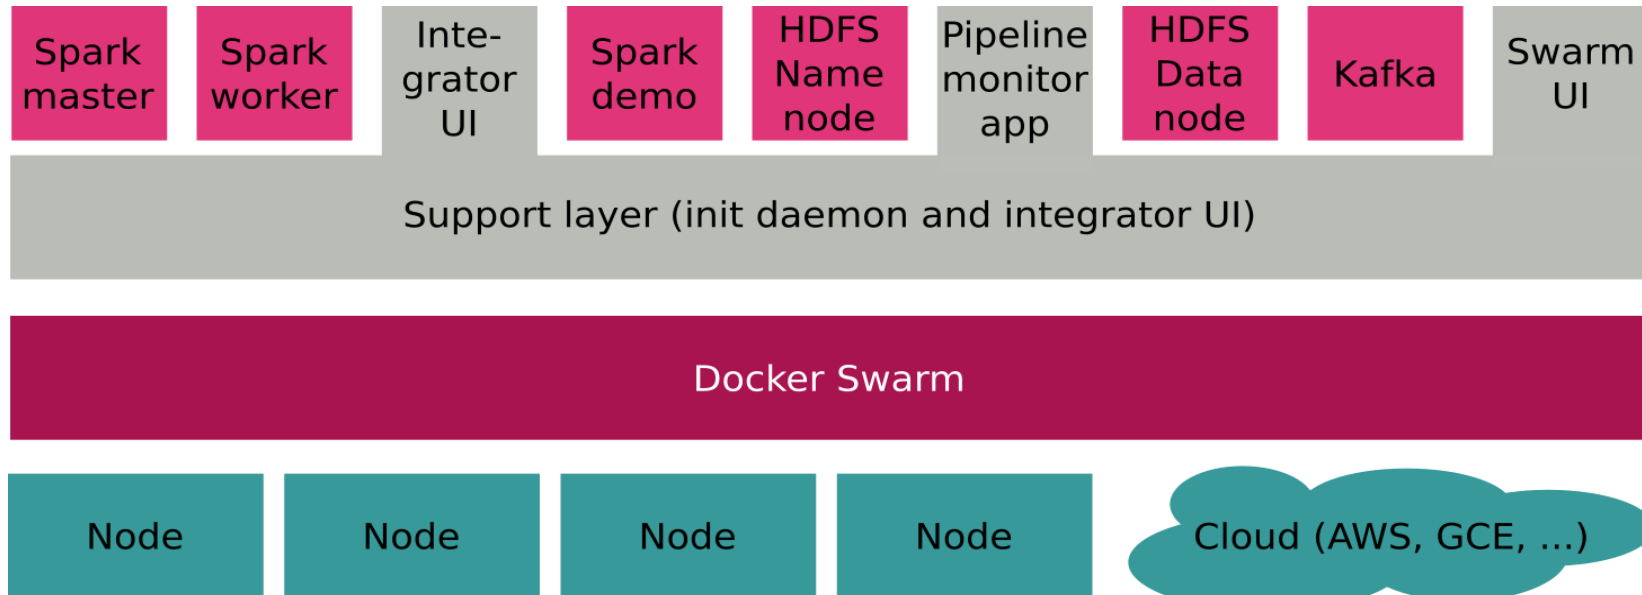
13





# Platform Architecture

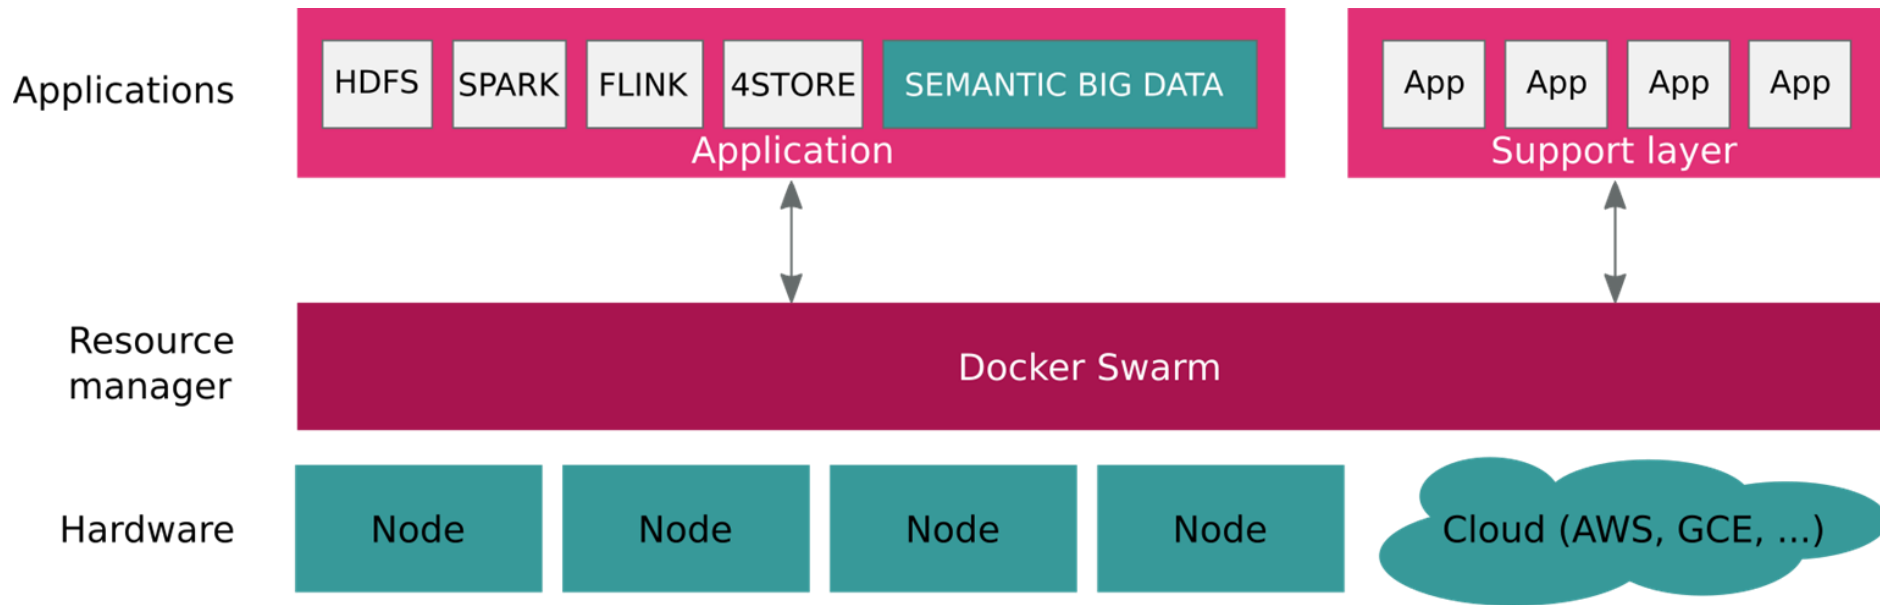
14





# Platform Architecture

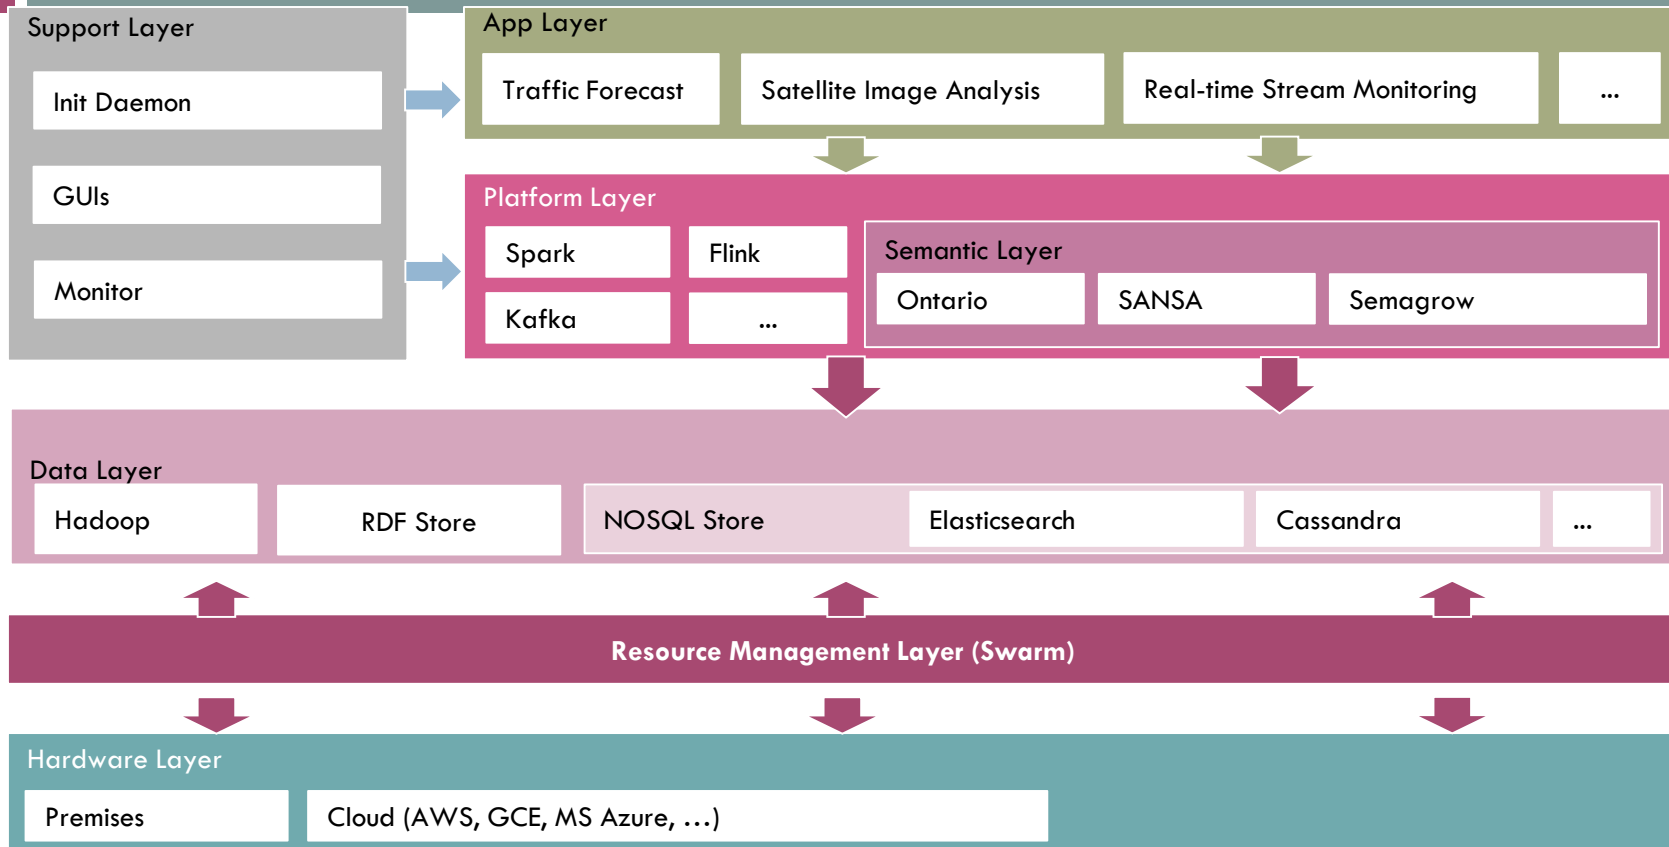
15





# Platform Architecture

16







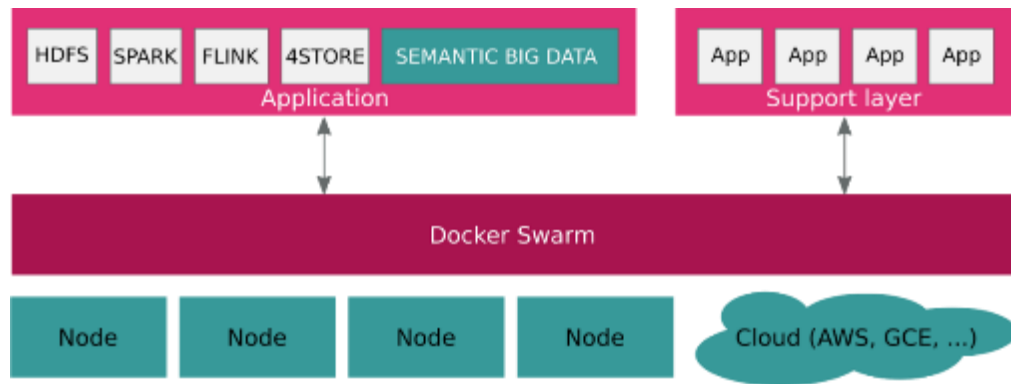
# BDI components

17

- ⊙ Processing and storage components
  - Re-used existing Docker containers where available
  - Dockerized by BDE where not
  - Ensured all can be provisioned through Docker Swarm

- ⊙ Components by BDE:

- Support Layer
- Semantic Layer





# BDE Docker Containers

18



Flink



SANSA





# Supported Frameworks

19

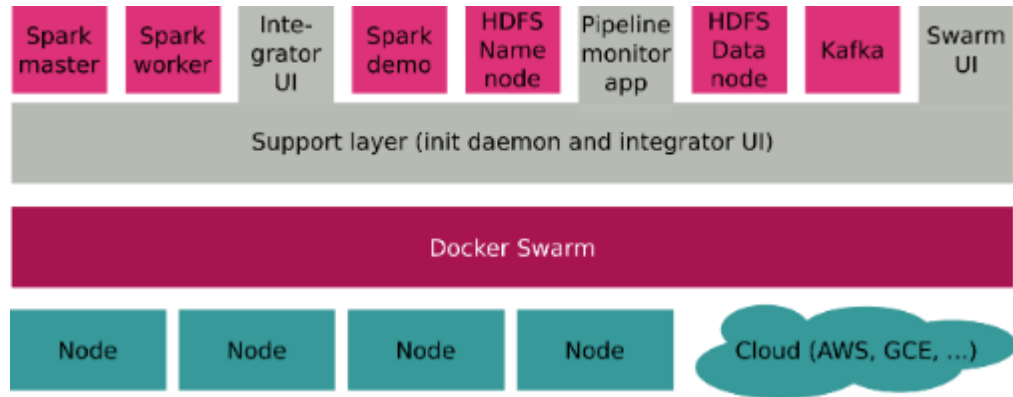
<b>Search/indexing</b>	<b>Data processing</b>
Apache Solr	Apache Spark
<b>Data acquisition</b>	Apache Flink
Apache Flume	<b>Semantic Components</b>
<b>Message passing</b>	Strabon
Apache Kafka	Sextant
<b>Data storage</b>	GeoTriples
Hue	Silk
Apache Cassandra	SEMAGROW
ScyllaDB	LIMES
Apache Hive	4Store
Postgis	OpenLink Virtuoso



# Support Layer

20

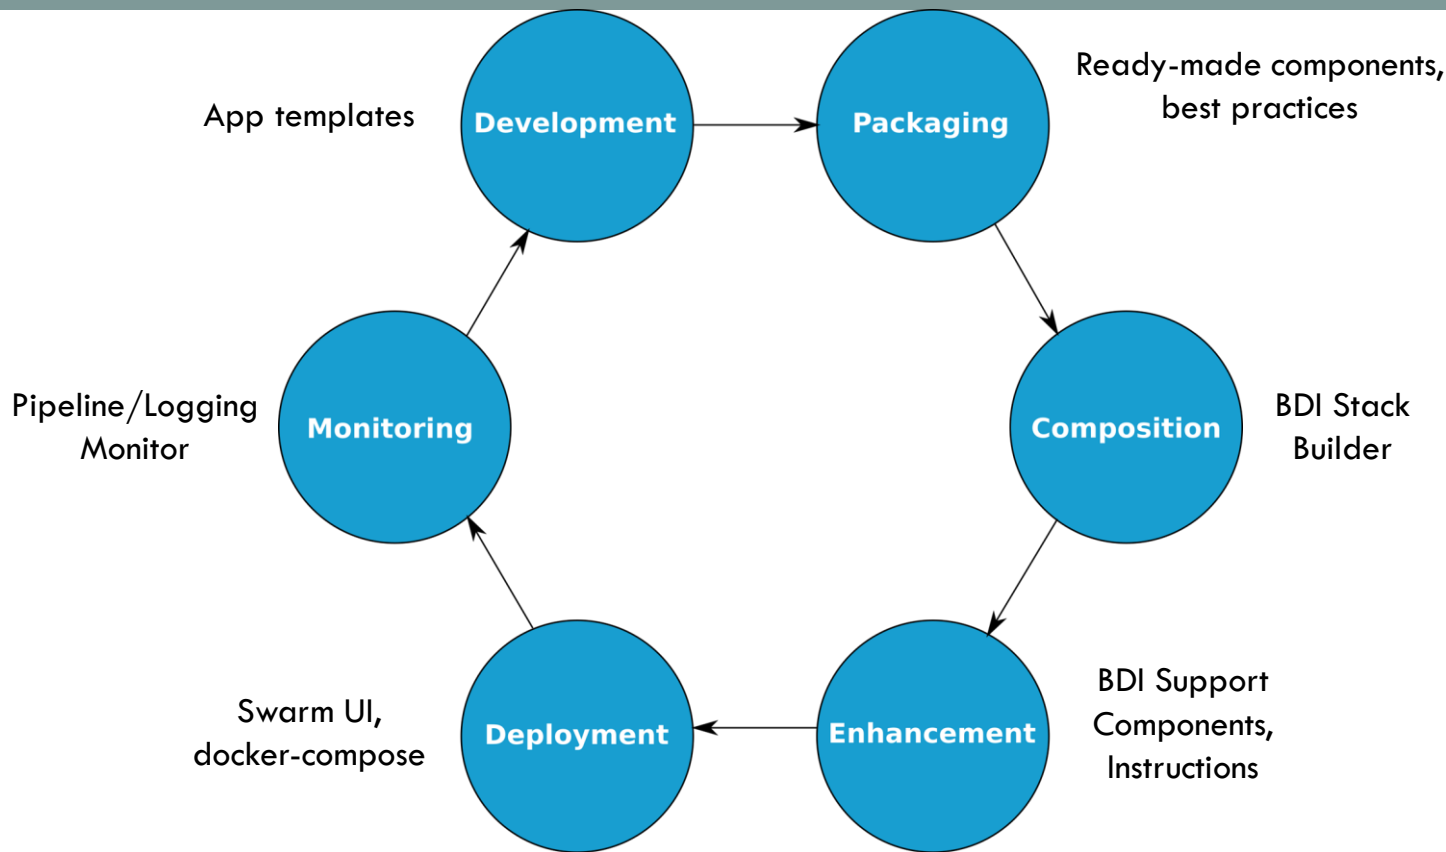
- ◎ BDE defines uniform UI stylesheets
  - Web UIs from BDE dockers (including for third party components) follow these BDE stylesheets
- ◎ BDE-developed tools:
  - Starting containers and dependencies
  - Monitoring execution





# BDI Stack Lifecycle

21





# BDI Stack Development

## ◎ High level picture

- docker-compose.yml describes pipeline topology

## ◎ BDE provided components

- extend template image with your code

## ◎ New components

- build a Docker image for your component
- this is your own little Virtual Machine for your component

## ◎ Sharing

- publish topology as git repository
- publish new components on docker hub




# Development

## ◎ Base Docker images

- Serve as a template for a (Big Data) technology
- Easily extendable custom algorithm/data

## ◎ Published components

- Image repositories on GitHub
- Automated builds on DockerHub
- Documentation on BDE Wiki





# Enhancing the Component

- ◎ Orchestrator required for initialization process (init\_daemon)
  - Components may depend on each other
  - Components may require manual intervention
- ◎ User Interface Integration
  - Standard Interfaces from components
  - Combine and align the interfaces





# User Interfaces

- ⊙ Target: Facilitate use of the platform
  - User Interface Adaption
- ⊙ Available interfaces
  - Workflow Uls
    - ❖ Workflow Builder
    - ❖ Workflow Monitor
  - Swarm UI
  - Integrator UI



# BDE Workflow Builder

26



BIG DATA EUROPE

BDE Workflow Builder

Workflows

## k-means demo

k-means Spark demo app

### Steps



#### Setup HDFS

Booting of the HDFS cluster.

setup\_hdfs

DELETE



#### Setup Spark

Starts the Spark master and workers.

setup\_spark

DELETE



#### Populate HDFS with core data

Please upload the location data to the HDFS filesystem. This is a manual step. Press finish when you're done

populate\_hdfs

DELETE



# BDE Workflow Monitor

27



BIG DATA EUROPE

BDE Workflow Monitor

Workflows

## Sensor demo

Vincent's fantastic sensor Spark app

### Steps



#### Setup HDFS

Booting of the HDFS cluster.

setup\_hdfs

FINISH



#### Setup SPARK

Starts the spark manager and workers.

setup\_spark

FINISH



#### Populate HDFS with core data

Please upload the location data to the HDFS filesystem. This is a manual step. Press finish when you're done

populate\_hdfs

FINISH



# Swarm UI

28

## Swarm UI

### Pipeline: WebCat

This pipeline is up.

Operations: UP STOP DOWN

#### Services

db

- 1 +

RESTART



resource

- 3 +

RESTART



dispatcher

- 3 +

RESTART



identifier

- 2 +

RESTART



webcat

- 1 +

RESTART





# Swarm UI

29



BIG DATA EUROPE

Swarm UI

## Pipeline: WebCat

This pipeline is up.

Operations:

UP

STOP

DOWN

## Services

identifier

-

1

+

RESTART



db

-

1

+

RESTART



2016-09-02T14:54:28.211935893Z

2016-09-02T14:54:28.211935893Z Fri Sep 02 2016

2016-09-02T14:54:28.212031995Z 14:54:28 { Loading plugin 1: Type 'plain', file 'wikiv' in '/usr/local/virtuoso-opensource/lib/virtuoso/hosting'

2016-09-02T14:54:28.212182746Z 14:54:28 FAILED plugin 1: Unable to locate file }

2016-09-02T14:54:28.212281608Z 14:54:28 { Loading plugin 2: Type 'plain', file 'mediawiki' in '/usr/local/virtuoso-opensource/lib/virtuoso/hosting'

2016-09-02T14:54:28.212390176Z 14:54:28 FAILED plugin 2: Unable to locate file }

2016-09-02T14:54:28.212467712Z 14:54:28 { Loading plugin 3: Type 'plain', file 'creolewiki' in '/usr/local/virtuoso-opensource/lib/virtuoso/hosting'

2016-09-02T14:54:28.212566398Z 14:54:28 FAILED plugin 3: Unable to locate file }

2016-09-02T14:54:28.213878861Z 14:54:28 OpenLink Virtuoso Universal Server

2016-09-02T14:54:28.213959049Z 14:54:28 Version 07.20.3212-pthreads for Linux as of Mar 14 2016

2016-09-02T14:54:28.214036041Z 14:54:28 uses parts of OpenSSL, PCRE, Html Tidy

2016-09-02T14:54:28.218459952Z 14:54:28 Database version 3126

2016-09-02T14:54:28.223186102Z 14:54:28 SQL Optimizer enabled (max 1000 layouts)

2016-09-02T14:54:29.409995744Z 14:54:29 Compiler unit is timed at 0.000200 msec

resource

-

1

+

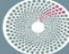
RESTART






# Integrator UI

30

BIG DATA EUROPE

Dashboard

MonitorVisualizationSpark MasterSpark WorkerHDFS NamenodeHueVirtuoso

1.6.2Spark Master at spark://spark-master:7077

URL: spark://spark-master:7077  
REST URL: spark://spark-master:6066 (cluster mode)  
Alive Workers: 1  
Cores in use: 3 Total, 0 Used  
Memory in use: 28.4 GB Total, 0.0 B Used  
Applications: 0 Running, 0 Completed  
Drivers: 0 Running, 0 Completed  
Status: ALIVE

Workers

Worker Id	Address	State	Cores	Memory
<a href="#">worker-20160901163503-172.18.0.13-33120</a>	172.18.0.13:33120	ALIVE	3 (0 Used)	28.4 GB (0.0 B Used)

Running Applications

Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
----------------	------	-------	-----------------	----------------	------	-------	----------

Completed Applications

Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
----------------	------	-------	-----------------	----------------	------	-------	----------



# Deploying a Big Data Stack

- ◎ Stack
  - collection of communicating components
  - to solve a specific problem
- ◎ Described in Docker Compose
  - Component configuration
  - Application topology



# Platform installation



Manual installation guide



Using Docker Machine

- On local machine (VirtualBox)
- In cloud (AWS, DigitalOcean, Azure)
- Bare metal



Screencasts





# Actors

33

- ◎ Cluster Setup
- ◎ Developer
- ◎ Packaging
- ◎ Stack Composition / Integration
- ◎ Deployment
- ◎ Monitoring



# BDE vs Hadoop distributions

34

	Hortonworks	Cloudera	MapR	Bigtop	BDE
<i>File System</i>	HDFS	HDFS	NFS	HDFS	<b>HDFS</b>
<i>Installation</i>	Native	Native	Native	Native	<b>lightweight virtualization</b>
<i>Plug &amp; play components (no rigid schema)</i>	no	no	no	no	<b>yes</b>
<i>High Availability</i>	Single failure recovery (yarn)	Single failure recovery (yarn)	Self healing, mult. failure rec.	Single failure recovery (yarn)	<b>Multiple Failure recovery</b>
<i>Cost</i>	Commercial	Commercial	Commercial	Free	<b>Free</b>
<i>Scaling</i>	Freemium	Freemium	Freemium	Free	<b>Free</b>
<i>Addition of custom components</i>	Not easy	No	No	No	<b>Yes</b>
<i>Integration testing</i>	yes	yes	yes	yes	--
<i>Operating systems</i>	Linux	Linux	Linux	Linux	<b>All</b>
<i>Management tool</i>	Ambari	Cloudera manager	MapR Control system	-	<b>Docker swarm UI+ Custom</b>



# BDE vs Hadoop distributions

- ⊙ BDE is not built on top of existing distributions
- ⊙ Consortium does not provide commercial support
- ⊙ Targets
  - Communities
  - Research institutions
- ⊙ Bridges scientists and open data
- ⊙ Multi Tier research efforts towards Smart Data



# Beyond the state of the art ...

## Smart Big Data

Increase the value of Big Data  
by adding meaning to it!



# Semantic layer tools

37

- ⊙ *Swagger*
- ⊙ *Semantic Analytics Stack (SANSA)*
- ⊙ *Semagrow*
- ⊙ *Ontario*
- ⊙ *LIMES*
- ⊙ *FOX*



**SemaGrow**



# Semantic Data Lake (Ontario)

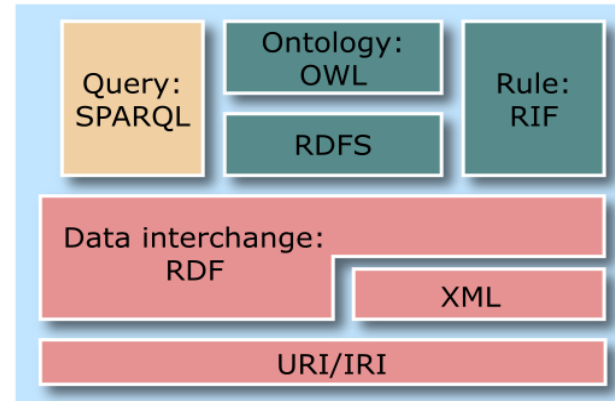

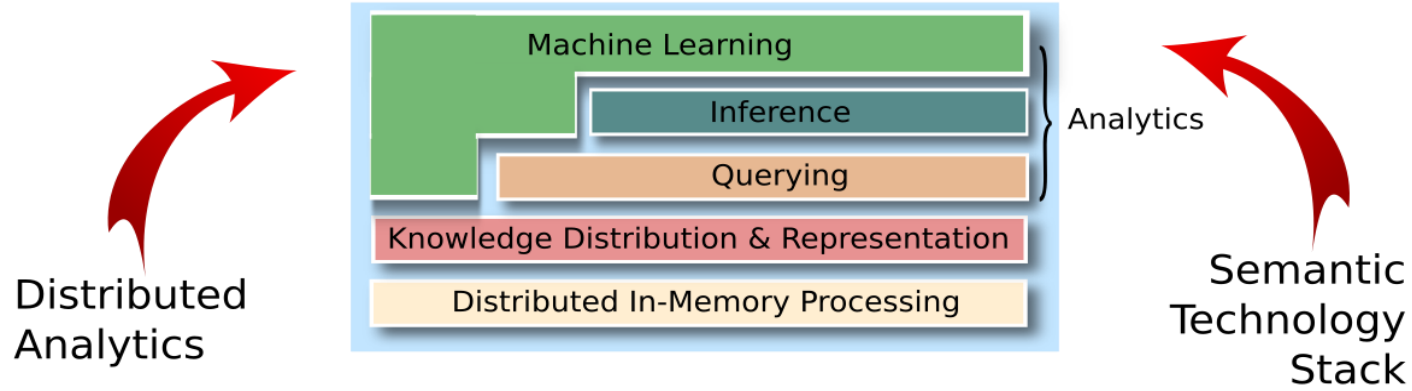
## ⊙ Data Swamp

- Repository of data in its raw format
- Structured, semi-structured, unstructured
- Schema-less

## ⊙ Data Lake

- Add a Semantic layer on top of the source datasets
- The data is semantically lifted using existing ontology terms

# Scalable Semantic Analytics Stack (SANSA)



- |                                       |                                       |
|---------------------------------------|---------------------------------------|
| - manual data integration             | + powerful data integration           |
| - often simple input formats          | + expressive modelling                |
| - data formats often not standardized | + W3C standardised formats            |
| + measurable benefits                 | - benefits only indirectly measurable |
| + horizontal scalability              | - usually no horizontal scalability   |



# SANSA Stack

40

Analytics

Classification

Link prediction

Clustering

Anomaly Detection

Latent Embeddings

...

Inference

Querying

Distribution

Knowledge Distribution and Representation

Distributed In-Memory Processing



41

# BigDataEurope Pilots




# SC1: Pharmacology research

42



Life  
Sciences  
& Health

- Query a large number of datasets, some large
- Existing elaborate ingestion and homogenization by the OpenPHACTS Foundation
- Extensive toolset developed by OPF and others





# SC1 Pilot: Points Demonstrated

43



Life  
Sciences  
& Health

- Existing distributed, scalable solution
  - Based on Virtuoso
  - Proprietary distributed triple store
- Porting to BDI gives flexibility
  - Using Virtuoso or open source alternatives (in BDE, 4store) without development effort for the superstructure and tools around it



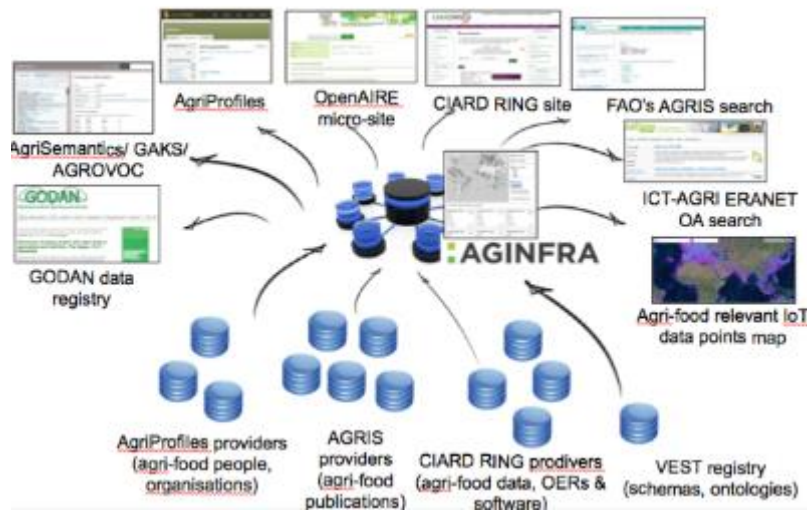
## SC2: Viticulture resources

44



Food and  
Agriculture

- AgInfra is a major infrastructure for agriculture researchers, serving cross-linked bibliography, data, and processing services
- Pilot automates ingestion and thematic classification of publication full texts





# SC2 Pilot: Points Demonstrated




## Food and Agriculture

- AglInfra: Existing infrastructure for data, metadata, and related services
- BDI is deployed as an external infrastructure for processing text (viticulture publications)
  - Allows storing and processing text at a larger scale than AglInfra can currently manage
  - The bibliographic metadata is added to AglInfra



# SC3: Predictive maintenance

46



Energy


- Wind turbine monitoring applies computational models to sensor data streams
- Models are weekly re-parameterized using week's data from multiple turbines





# SC3 Pilot: Points Demonstrated

47



Energy

- Existing in-house non-scalable solution for model parameterization
  - Reliable Fortran software for data analysis
  - Efficient, but not scalable to data volume
- Developing a BDI orchestrator
  - Re-uses existing software unmodified
  - Makes it easy to apply in parallel to many datasets and manage the outputs



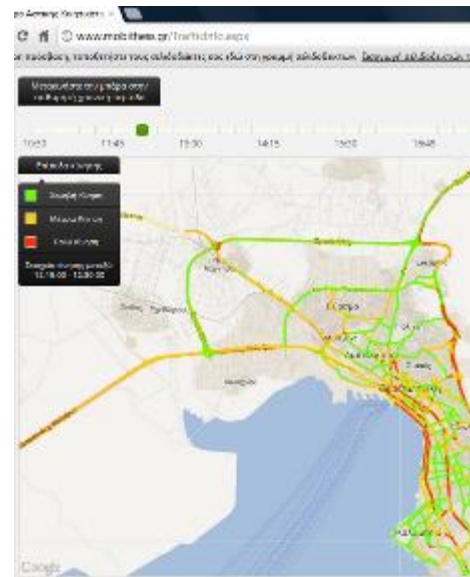
# SC4: Traffic conditions estimation

48



Transport

- Estimation of real-time traffic conditions in Thessaloniki
- Combines:
  - Traffic modelling from historical data
  - Current measurements from a taxi fleet of 1200 vehicles







# SC4 Pilot: Points Demonstrated

49



Transport

- New Flink implementations of map matching and traffic prediction algorithms
- BDI provides access to varied data sources
  - PostGIS database with city map
  - ElasticSearch database of historical data
  - Kafka stream of real-time data




# SC5: Climate modelling

50



Climate

- Preparing modelling experiments
  - Slicing, transforming, combining datasets into new datasets
  - Submission to and retrieval from modelling infrastructure
- Discovering and re-using previously computed derivatives
  - Lineage annotation: datasets and model parameters used to compute derivative datasets
  - Finding appropriate past runs avoids repeating weeks-long modelling runs





# SC5 Pilot: Points Demonstrated

51



Climate

- Existing infrastructure and stable, reliable software for parallel computation of models
- BDI is deployed as an external infrastructure for preparing and managing datasets
- BDI offers:
  - Hive for managing data in a way that can be retrieved and manipulated, rather than file blocks
  - Cassandra stores structured and textual metadata for searching headers and lineage



# SC6: Municipality budgets

52



Social  
Sciences

- Ingestion of budget and budget execution data
- Multiple municipalities in varied formats and data models
- Homogenized data made available for analysis and comparison





# SC6 Pilot: Points Demonstrated



## Social Sciences

- Existing analytics and visualization tools
  - Use SPARQL queries to retrieve only the relevant slices of the overall data
- BDI is deployed as an ingestion and storage infrastructure for external tools
  - Ingests and homogenizes a constant flow of JSON, CSV, XML, and other formats following various data models
  - Exposes data as SPARQL endpoint serving homogenized data, stored in 4store, a scalable, distributed RDF store




# SC7: Change detection & verification

54



Secure  
Societies

- Events are extracted from text published by news agencies and on social networking sites
- Events are geo-located and relevant changes are detected by comparing current and previous satellite images





# SC7 Pilot: Points Demonstrated

55



Secure  
Societies

- Re-implementation of change detection algorithms for Spark
- Parallel orchestrator for text analytics
  - Re-uses existing software
  - Scales to many input streams
- BDI provides:
  - Cassandra for text content and metadata
  - Strabon GIS store for detected change location
  - Homogeneous access to both for analysis and visualization

56

# Closing Remarks





# BDI use cases demonstrated

## ⊙ *Flexibility to existing workflows*

- Drug discovery platform ported to new data storage infrastructure

## ⊙ *Scaling out*

- Sensor data analysis software trivially distributed to multiple streams
- Image analysis and traffic pattern algorithms in Spark and Flink

## ⊙ *Preprocessing and ingestion*

- Distributed data architectures pre-process and maintain big data
- Not needed simultaneously or reduced to small data



# Semantic Web and Big Data

- ⊙ *Adding triple stores to the Apache ecosystem*
  - Scalability for Semantic Web technologies
  - Semantics stores for the Apache ecosystem
- ⊙ *Advocating RDF technologies for describing processing and data services*
  - The Semantic Data Lake vision builds an RDF layer over Swagger/OAI
- ⊙ *Advocating RDF and SPARQL as the lingua franca for Big Data retrieval*
  - Semagrow federations of heterogeneous internal and external data stores



# Thank you for your attention! Questions?



**BIG DATA EUROPE**

Empowering Communities  
with Data Technologies

- ◎ BigDataEurope action Web site:  
<https://www.big-data-europe.eu>
- ◎ Big Data Integrator:  
<https://github.com/big-data-europe>
- ◎ Semagrow:  
<http://semagrow.github.io>
- ◎ SANSA:  
<http://sansa-stack.net>