BIG DATA EUROPE

Empowering Communities
with Data Technologies

Planning BDI Stack for your Big Data Application

Ivan Ermilov @ ICTCS, Amman, Jordan

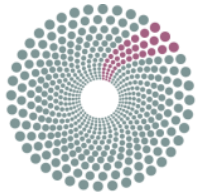

Outline

- Dockerization of Big Data Frameworks (BDF): Why
- What is BDI Stack?
- BDI Stack Lifecycle
- BDI Stack Assembly
- Examples
  - New Spark application
  - Existing Spark application

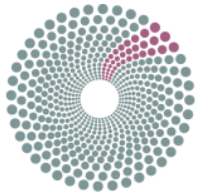

Dockerization of BDF: Why

	Development VM	QA Server	Single Prod Server	Onsite Cluster	Public Cloud	Contributor's laptop	Customer Servers
Static website	?	?	?	?	?	?	?
Web frontend	?	?	?	?	?	?	?
Background workers	?	?	?	?	?	?	?
User DB	?	?	?	?	?	?	?
Analytics DB	?	?	?	?	?	?	?
Queue	?	?	?	?	?	?	?

docker

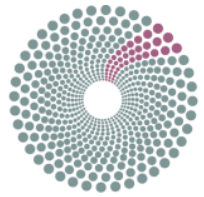

Dockerization of BDF: Why

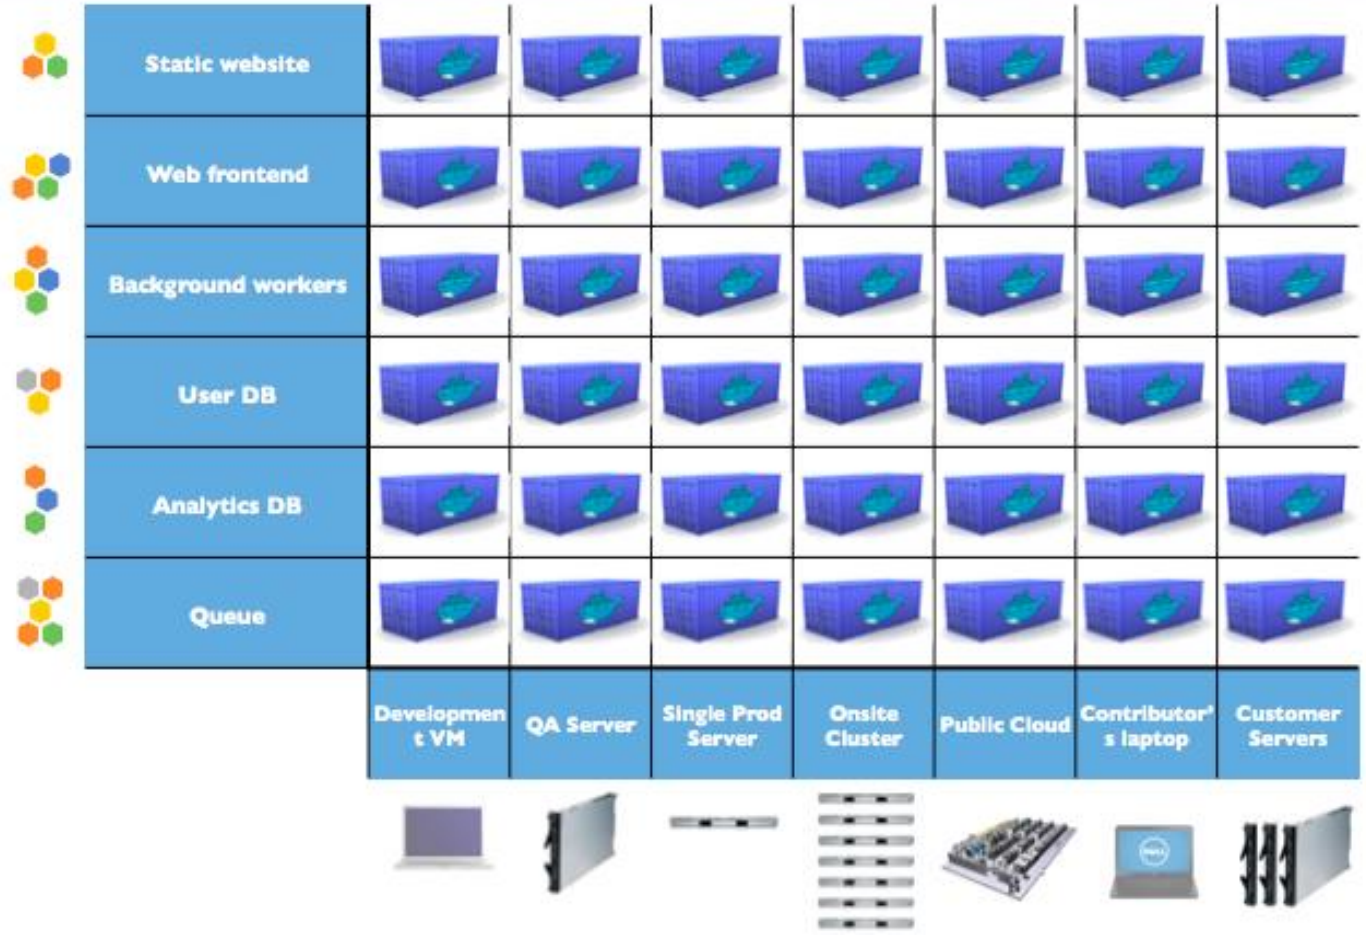

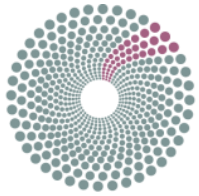

Dockerization of BDF: Why

- Development environment
- Testing environment
- Staging environment
- Production environment

They all the same!

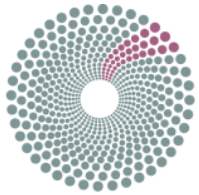

Less Duplication = Less Bugs

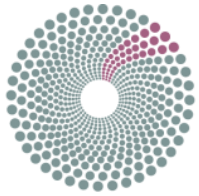

What is BDI Stack?

- Dockerized BDF
- In one bundle
- With custom applications
- docker-compose.yml

BDI Stack Lifecycle

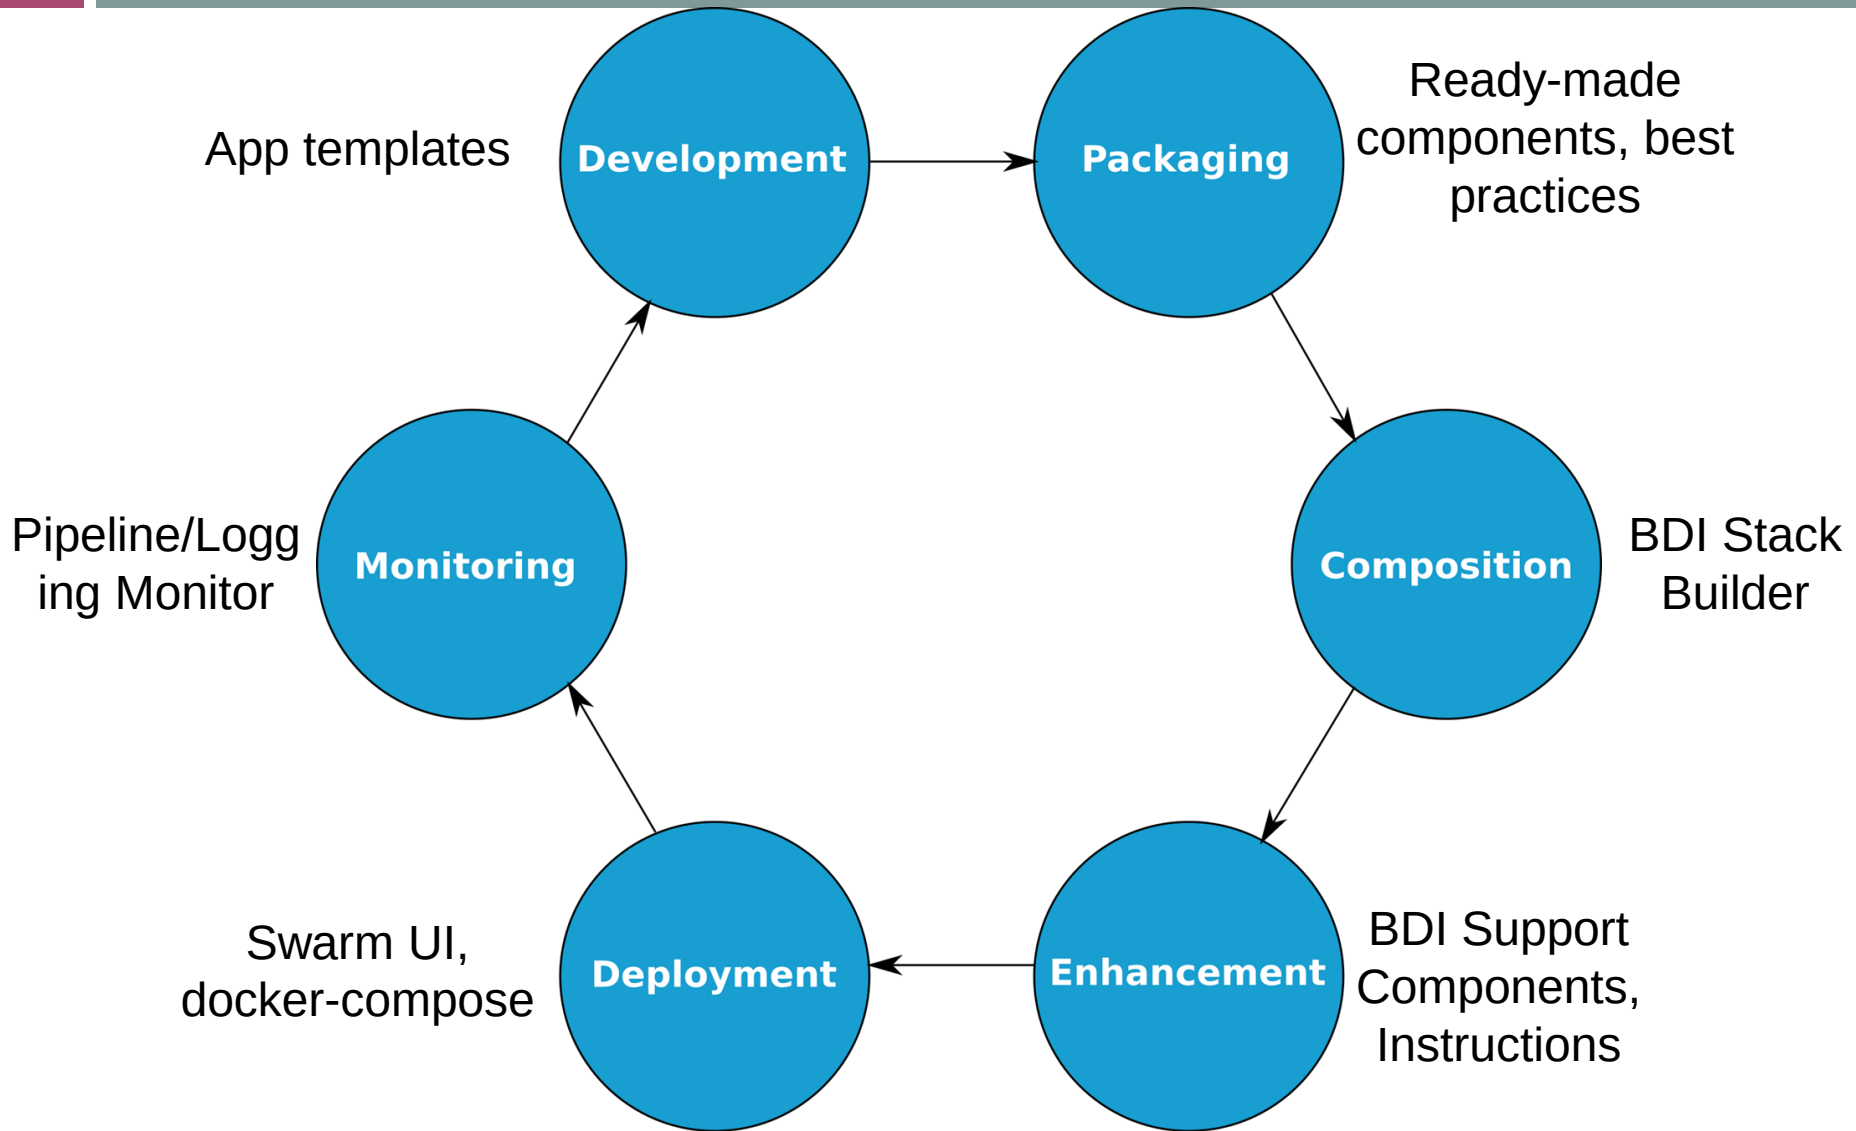

BDI Stack Assembly

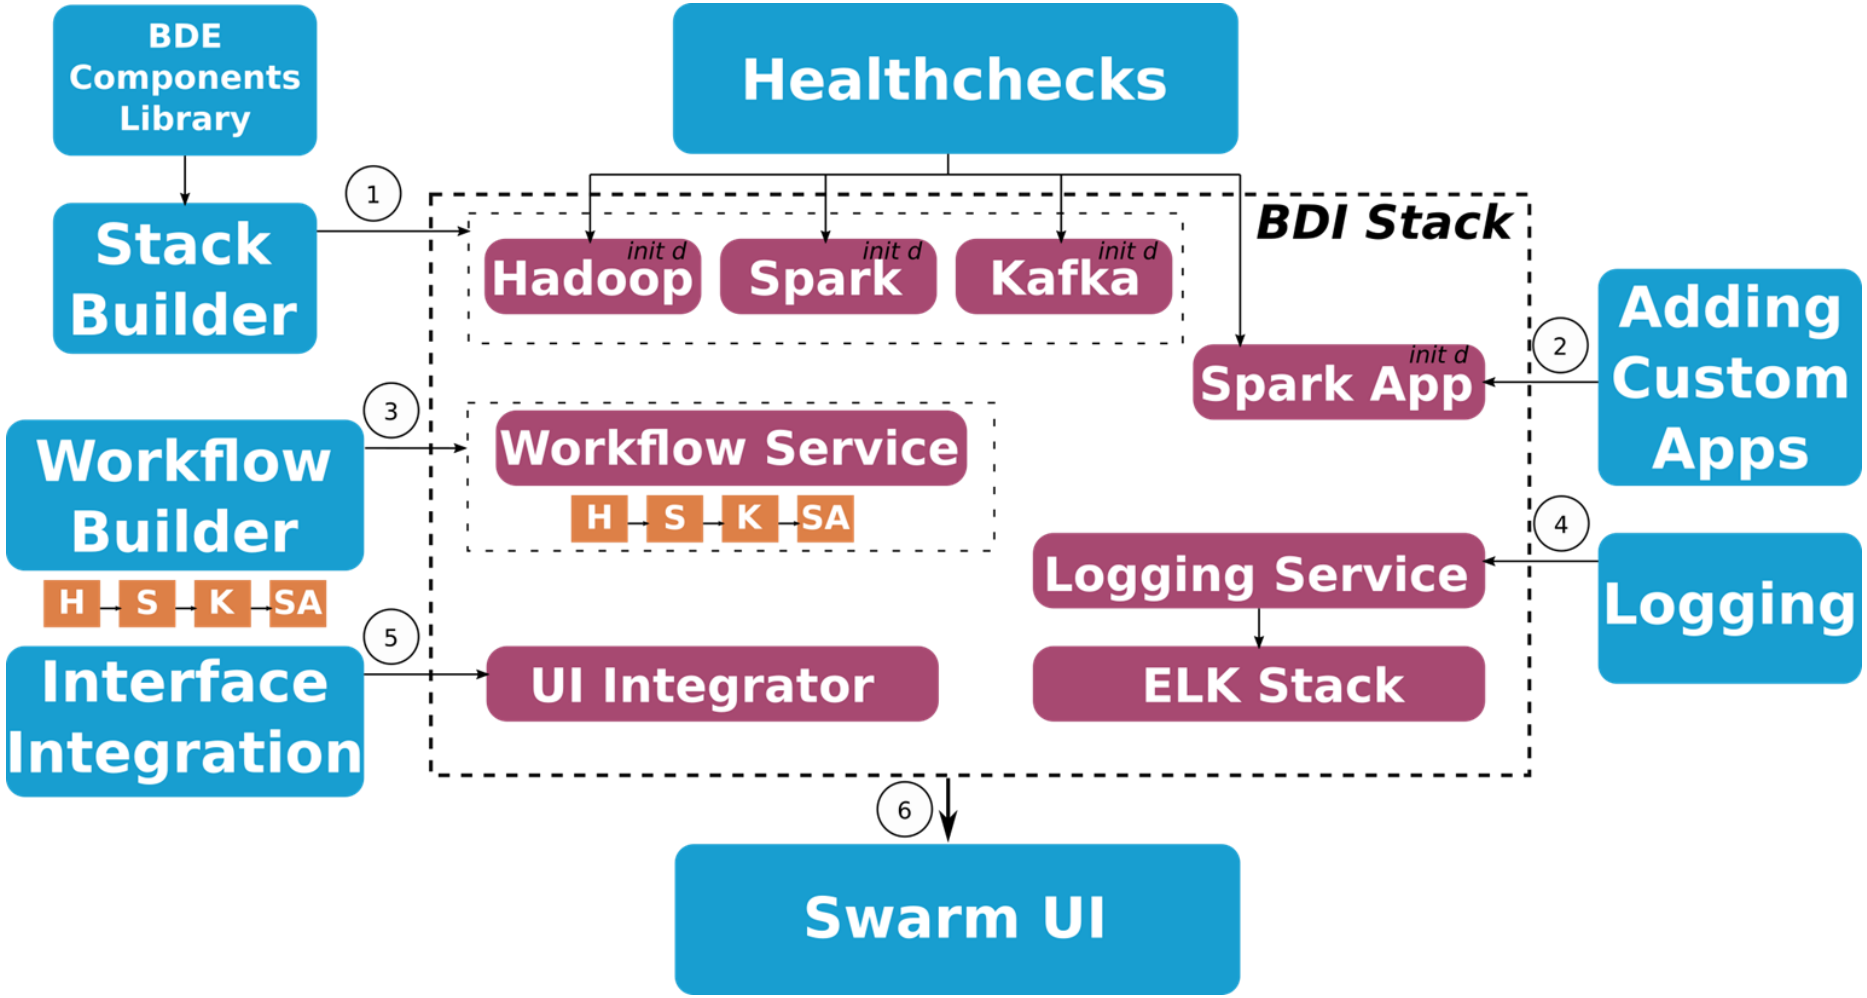

Stack Builder

Stack Builder

Compact list view

Title

docker-spark

Text

```
version: "2"
services:
  spark-master:
    image: bde2020/spark-master:2.1.0-hadoop2.7
    container_name: spark-master
    ports:
      - "8080:8080"
      - "7077:7077"
    environment:
      - INIT_DAEMON_STEP=setup_spark
      - "constraint:node==<yourmasternode>"
  spark-worker-1:
    image: bde2020/spark-worker:2.1.0-hadoop2.7
    container_name: spark-worker-1
    depends_on:
      - spark-master
    ports:
      - "8081:8081"
    environment:
```

Search

To add items from the list to your docker-compose file simply drag and drop them into the text field.

bde2020/4store	
bde2020/4store-master	
flink	2
hdfs	6

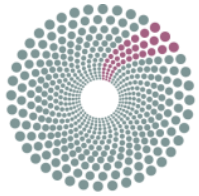

Adding Custom Apps


```
FROM bde2020/spark-submit:2.1.0-hadoop2.7

ENV ENABLE_INIT_DAEMON=false

ENV SPARK_APPLICATION_PYTHON_LOCATION=

ENV SPARK_MASTER_NAME=sc6-spark-master

ENV SPARK_APPLICATION_ARGS=

ENV SPARK_MASTER_URL=spark://sc6-spark-master:7077

ENV SPARK_MASTER_PORT=7077
```

WorkFlow Builder

BDE Workflow Builder

Workflows

My Workflow

My Workflow

Steps

init-hadoop

Init Hadoop

Initialization of Hadoop

DELETE

init-spark

Init Spark

Initialization of Spark

DELETE

step_code

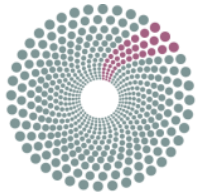

Logging Monitor

- Network logging for HTTP
  - Capture network interface as PCAP
  - Convert to HAR (json)
  - Expand HAR
  - Dump into ELK stack

UI Integrator Application

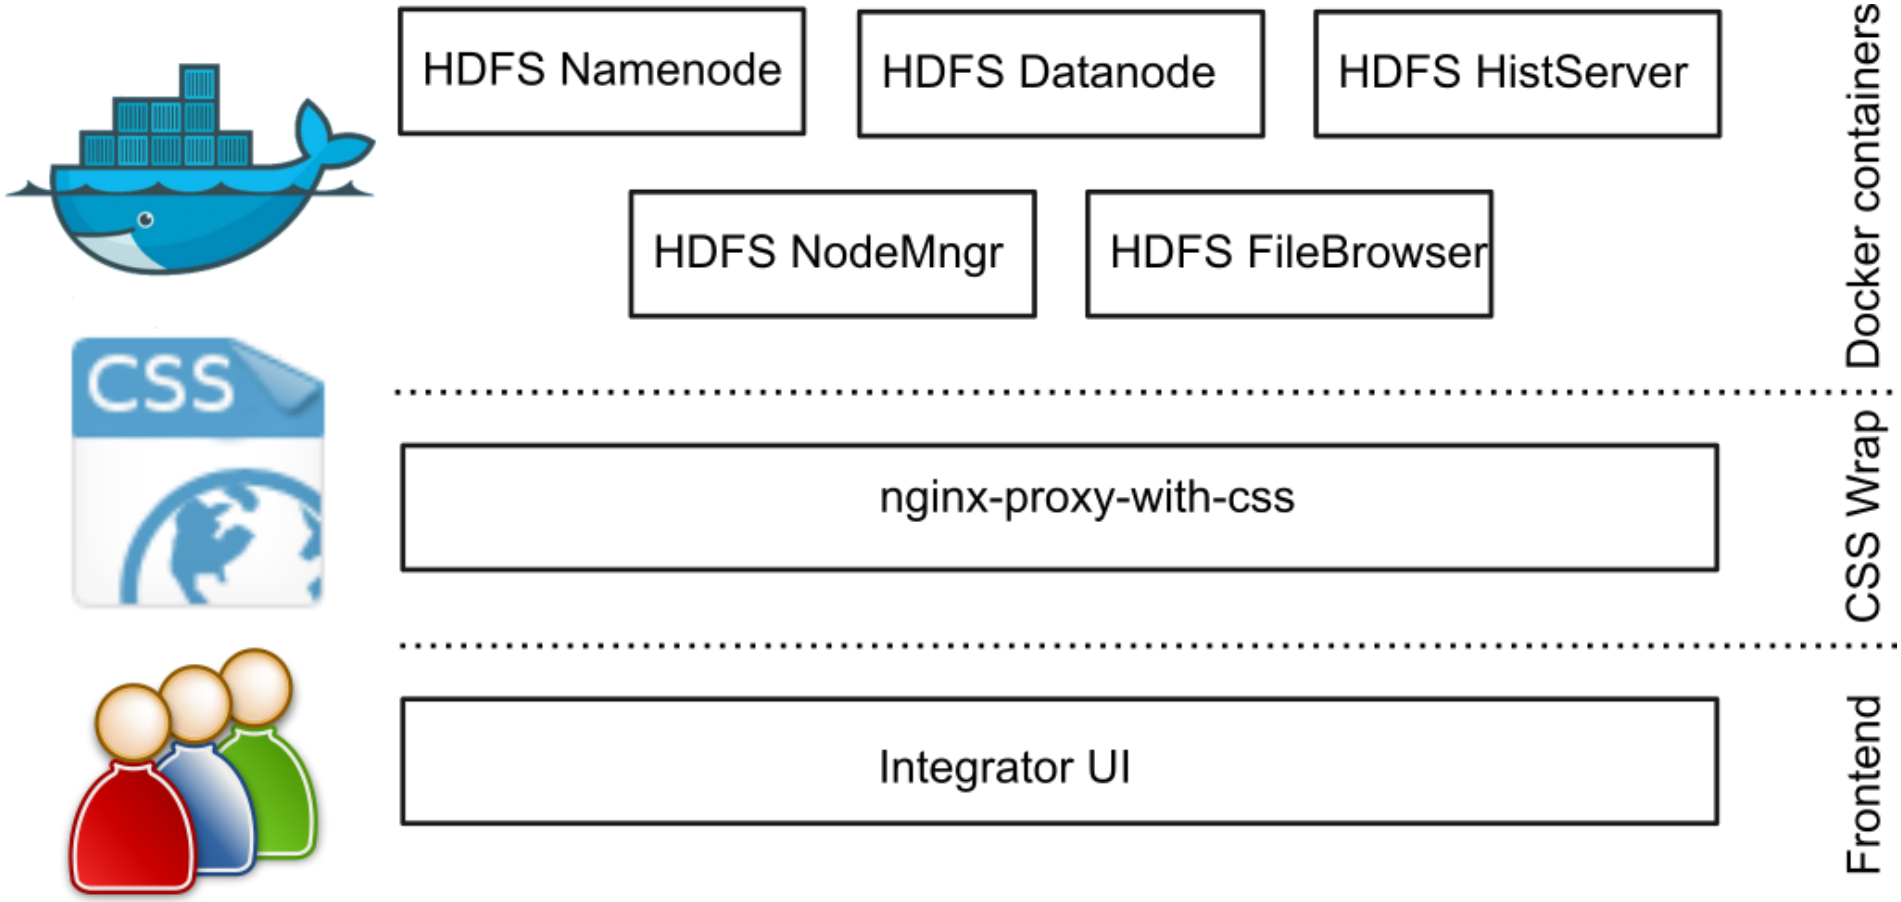

https://www.big-data-europe.eu/user-interface-integration-in-bdi-platform-integrator-ui-application/

UI Integrator

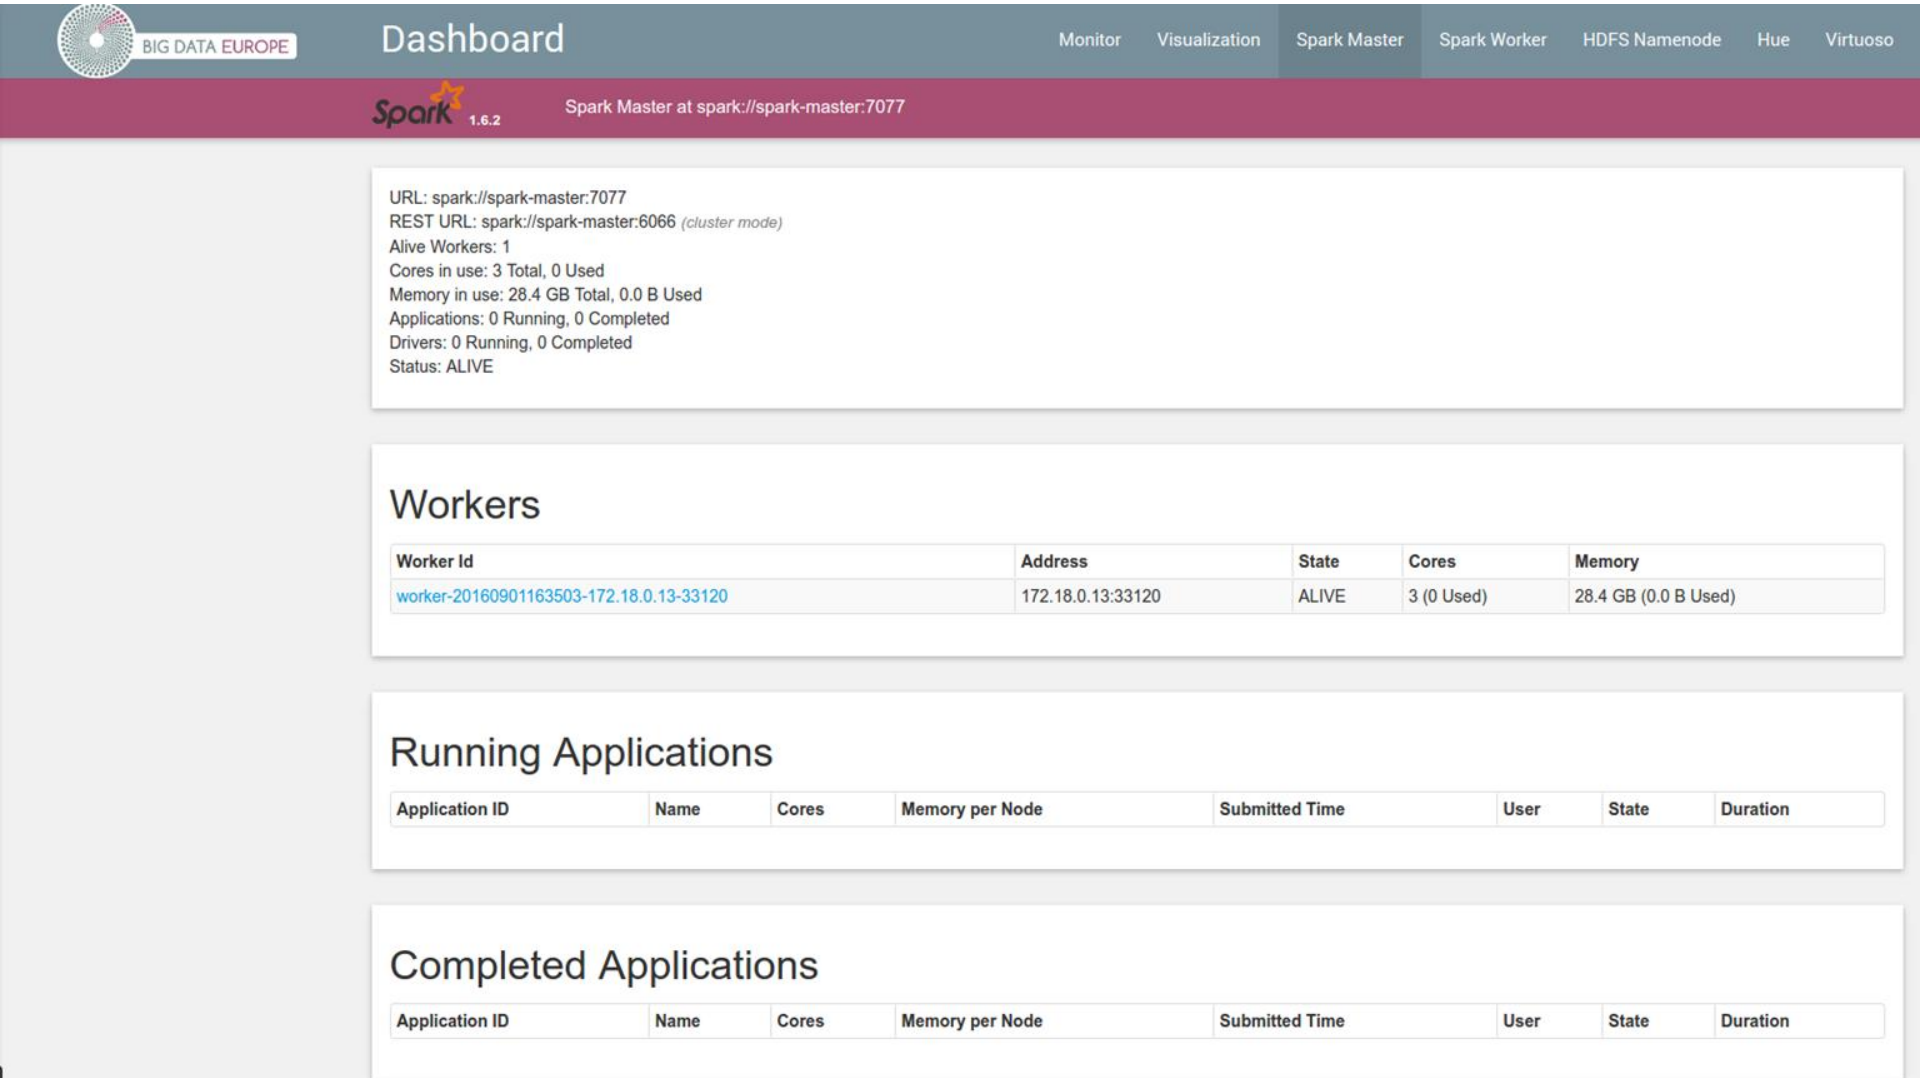

Reverse Proxy/CSSWrapper

- Simple injection of custom CSS

```
strabon:
  image: bde2020/strabon
  links:
    - csswrapper
  expose:
    - "8080"
  environment:
    VIRTUAL_HOST: "strabon.big-data-europe.aksw.org"
    VIRTUAL_PORT: "8080"
    CSS_SOURCE: "strabon"
```

https://www.big-data-europe.eu/using-reverse-proxy-inside-bde-platform-jwilderngінx-setup-for-docker-swarm/
https://github.com/big-data-europe/demo-integrator-ui

Swarm UI

Swarm UI

Repositories Pipelines

Repositories

test

Located at https://github.com/big-data-europe/demo-spark-sensor-data. Has 1 connected pipelines.

EDIT LAUNCH

Create new repository

Swarm UI

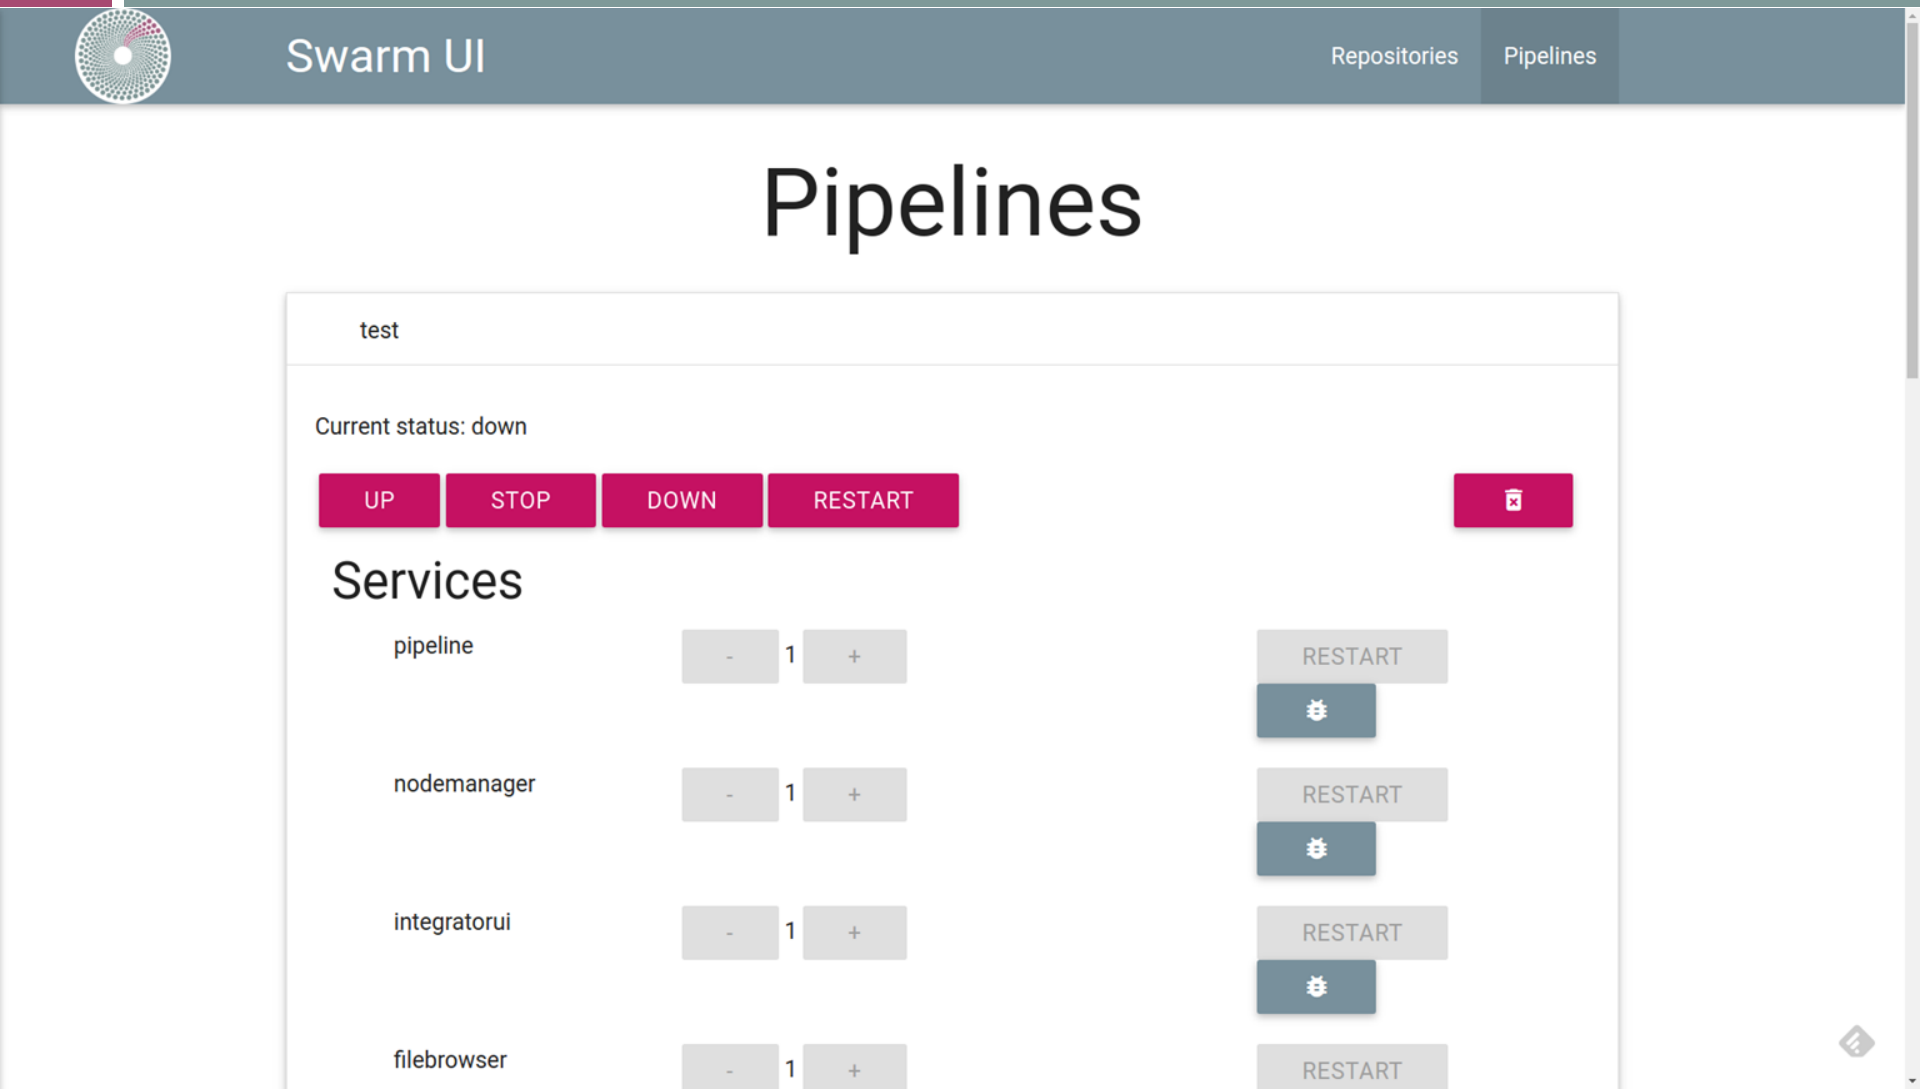
Swarm UI
Repositories
Pipelines
Pipelines
test
Current status: down
UP
STOP
DOWN
RESTART
Services
pipeline
1
RESTART
nodemanager
1
RESTART
integratorui
1
RESTART
filebrowser
1
RESTART

BDI Stack Assembly

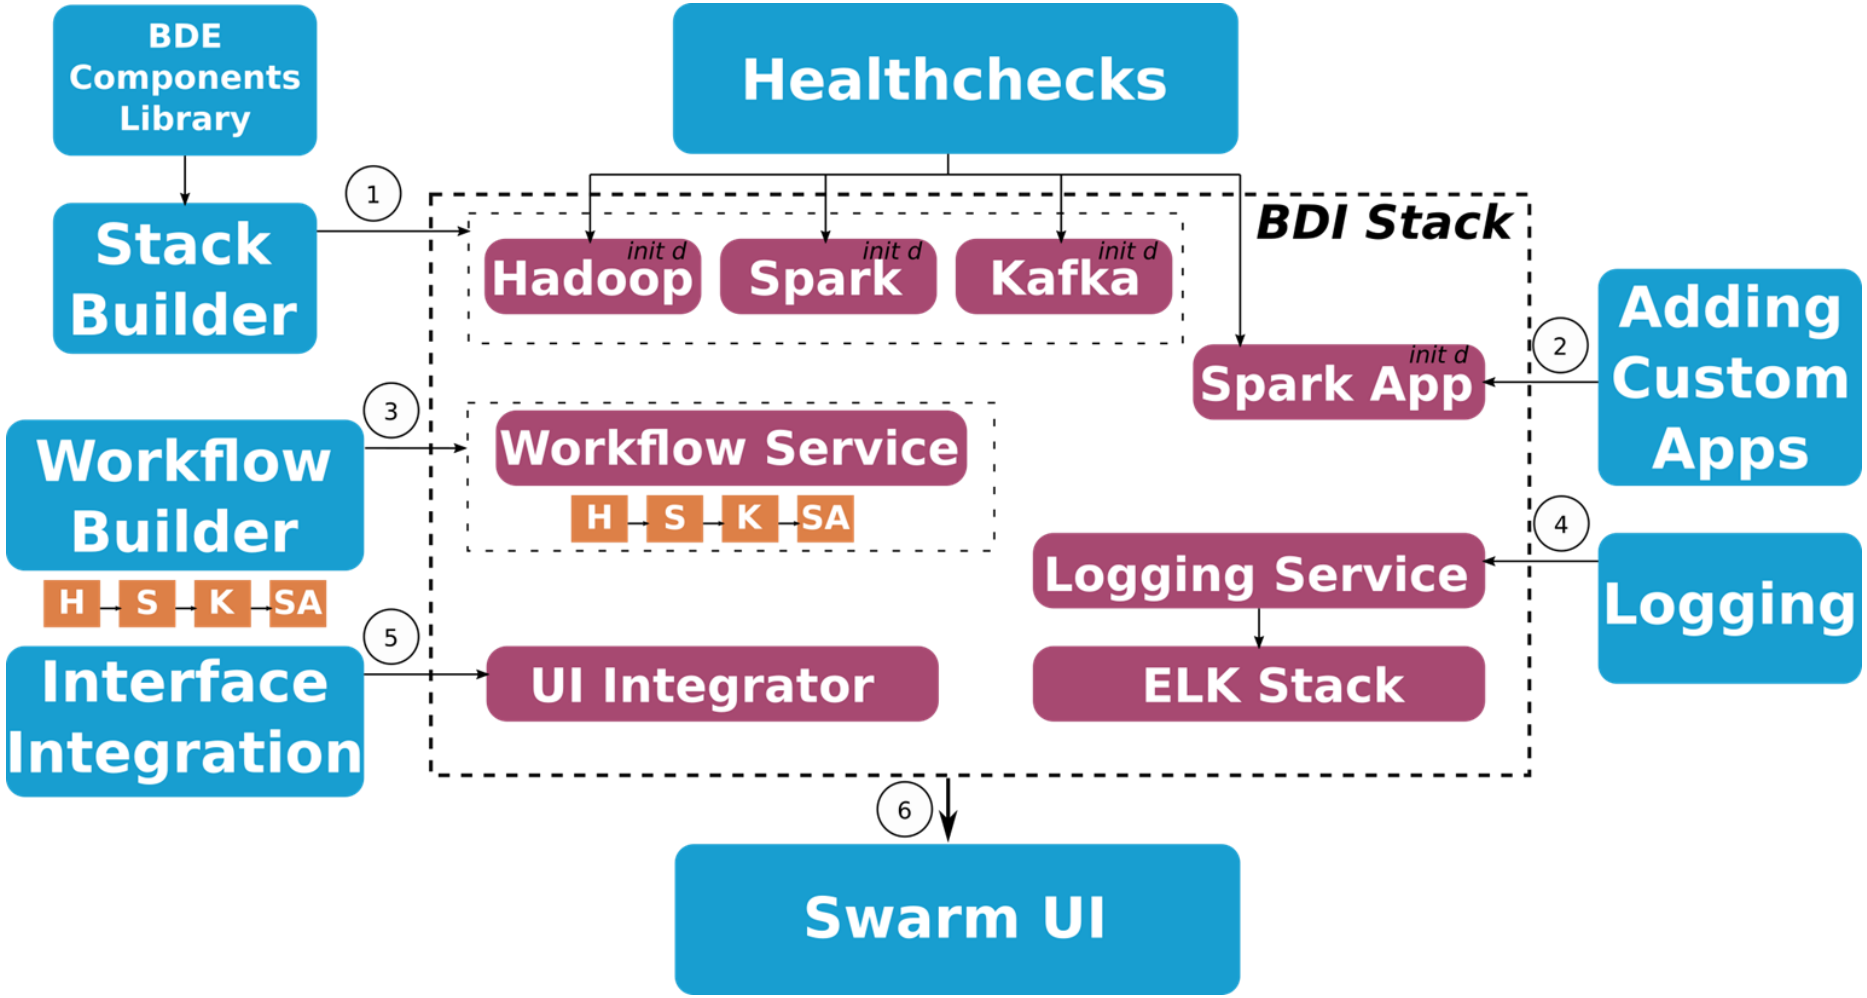

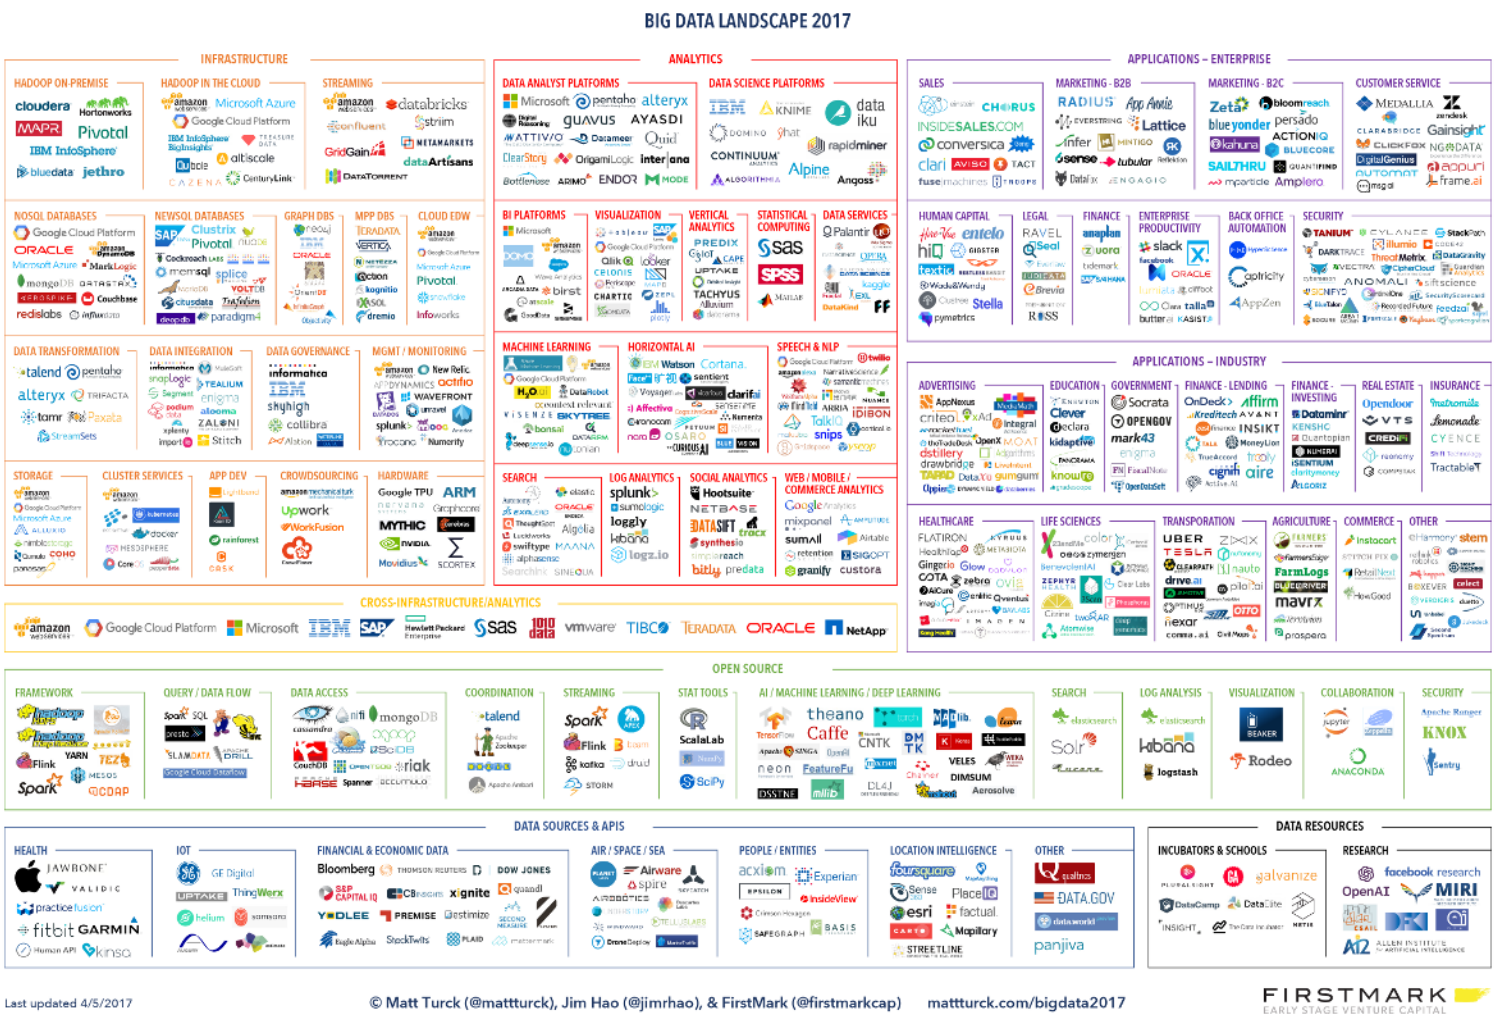
BIG DATA LANDSCAPE 2017
INFRASTRUCTURE
HADOOP ON-PREMISE
HADOOP IN THE CLOUD
STREAMING
NOSQL DATABASES
NEWSQL DATABASES
GRAPH DBS
MPP DBS
CLOUD EDW
DATA TRANSFORMATION
DATA INTEGRATION
DATA GOVERNANCE
MGMT / MONITORING
STORAGE
CLUSTER SERVICES
APP DEV
CROWDSOURCING
HARDWARE
ANALYTICS
DATA ANALYST PLATFORMS
DATA SCIENCE PLATFORMS
BI PLATFORMS
VISUALIZATION
VERTICAL ANALYTICS
STATISTICAL COMPUTING
DATA SERVICES
MACHINE LEARNING
HORIZONTAL AI
SPEECH & NLP
SEARCH
LOG ANALYTICS
SOCIAL ANALYTICS
WEB / MOBILE / COMMERCE ANALYTICS
CROSS-INFRASTRUCTURE/ANALYTICS
APPLICATIONS – ENTERPRISE
SALES
MARKETING - B2B
MARKETING - B2C
CUSTOMER SERVICE
HUMAN CAPITAL
LEGAL
FINANCE
ENTERPRISE PRODUCTIVITY
BACK OFFICE AUTOMATION
SECURITY
APPLICATIONS – INDUSTRY
ADVERTISING
EDUCATION
GOVERNMENT
FINANCE - LENDING
FINANCE - INVESTING
REAL ESTATE
INSURANCE
HEALTHCARE
LIFE SCIENCES
TRANSPORTATION
AGRICULTURE
COMMERCE
OTHER
OPEN SOURCE
FRAMEWORK
QUERY / DATA FLOW
DATA ACCESS
COORDINATION
STREAMING
STAT TOOLS
AI / MACHINE LEARNING / DEEP LEARNING
SEARCH
LOG ANALYSIS
VISUALIZATION
COLLABORATION
SECURITY
DATA SOURCES & APIS
HEALTH
IOT
FINANCIAL & ECONOMIC DATA
AIR / SPACE / SEA
PEOPLE / ENTITIES
LOCATION INTELLIGENCE
OTHER
DATA RESOURCES
INCUBATORS & SCHOOLS
RESEARCH
Last updated 4/5/2017
© Matt Turck (@mattturck), Jim Hao (@jimrhao), & FirstMark (@firstmarkcap)
mattturck.com/bigdata2017
FIRSTMARK
EARLY STAGE VENTURE CAPITAL

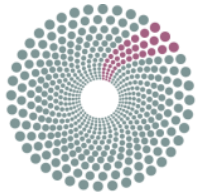

Simplified Workflow

- Outline your application requirements
- Pick up components from BDE github repo
  - If it is not there search Docker Hub
  - Else dockerize it yourself
- Create docker-compose.yml
- Test with simple application
- Develop your application on top of it
- Proceed to enhancement step (if necessary)

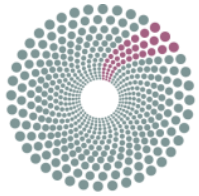

Application Requirements

- Create user stories
  - User wants to see the most trending recent hashtags
  - User wants to see the most recent visualization of the hashtags
  - User wants to see past visualizations as well

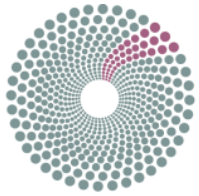

Core Functionality

- Fetch tweets
  - Spark streaming can do that
- Store tweets somewhere (big data)
  - HDFS
- Store trends (not too big)
  - MongoDB
- Visualize the trends
  - Kibana or custom application

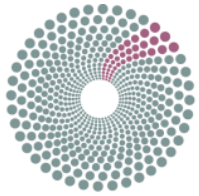

Components from BDE github

- HDFS
  - Store data
- Spark
  - Streaming
  - Transformation
  - Save/load data

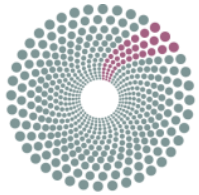

Searching Docker Hub

Search - Docker Hub

hub.docker.com/search/?isAutomated=0&isOfficial=0&page=1&pullCount=0&q=kafka&starCount=0

Docker Store is the new place to discover public Docker content. Check it out →

kafka

Explore Help Sign up Sign in

Repositories (3097)

All

wurstmeister/kafka public | automated build	417 STARS	10M+ PULLS	DETAILS
sheepkiller/kafka-manager public | automated build	80 STARS	1M+ PULLS	DETAILS
ches/kafka public | automated build	82 STARS	1M+ PULLS	DETAILS
confluentinc/cp-kafka-connect public	12 STARS	1M+ PULLS	DETAILS
spotify/kafka public | automated build	225 STARS	500K+ PULLS	DETAILS
cgswong/confluent-kafka public | automated build	7 STARS	1M+ PULLS	DETAILS
solsson/kafka			

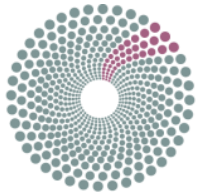

Searching Docker Hub

- Create a table (example for HBase)
  - Name
  - Java Version
  - Docs
  - Configurable
  - Standalone
  - Pseudodistributed
  - Distributed

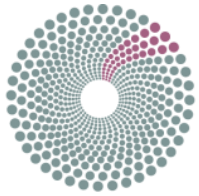

Searching Docker Hub

- Select the best docker image from the table
  - Pay attention to the docker image license!
- Extend if necessary
- Example

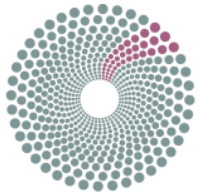

Assembling docker-compose

- BDE provides docker-compose.yml snippets
- Docker images which follow the best practices does the same

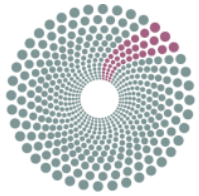

Assembling docker-compose

- Copy/paste the snippets and adjust

```
version: '2'
services:
  namenode:
    image: bde2020/hadoop-namenode:1.1.0-hadoop2.8-
java8
    container_name: namenode
    volumes:
      - ./data/namenode:/hadoop/dfs/name
    environment:
      - CLUSTER_NAME=TwitterTrendsCluster
…
```

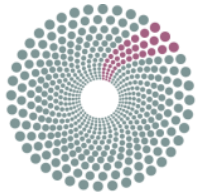

Testing your BDI Stack

- Manual testing
  - All containers running?
  - No errors in the initialization logs?
- Automatic testing
  - Deploy a simple application using the BDI stack along it
  - Produce correct results?

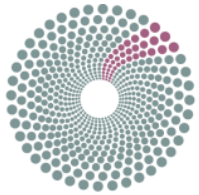

Ready to develop your app!

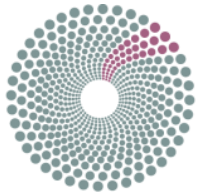

Adding Application to the Stack

- Create Dockerfile
- Expose external interfaces
  - REST
  - SQL
  - SPARQL
- Upload to docker hub
  - Or your enterprise repository (e.g. gitlab)

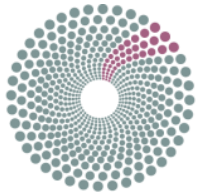

Demo (15 mins)

- BDI Stack for
  - Hadoop
  - Spark
  - TwitterTrends
  - VisualizationApp

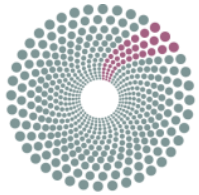

Packing existing application

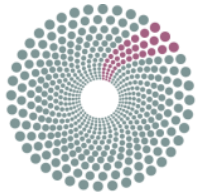

Example BDI Stack: Halyard

- Which BD components does Halyard use?
  - HDFS
  - YARN (for MapReduce jobs)
  - HBase
- Which interfaces are supported?
  - Shell scripts (bulkload, export etc)
  - RDF4J console
  - RDF4J REST Server + Workbench

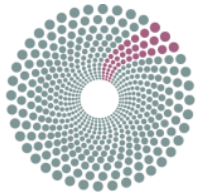

Halyard: BDI Stack

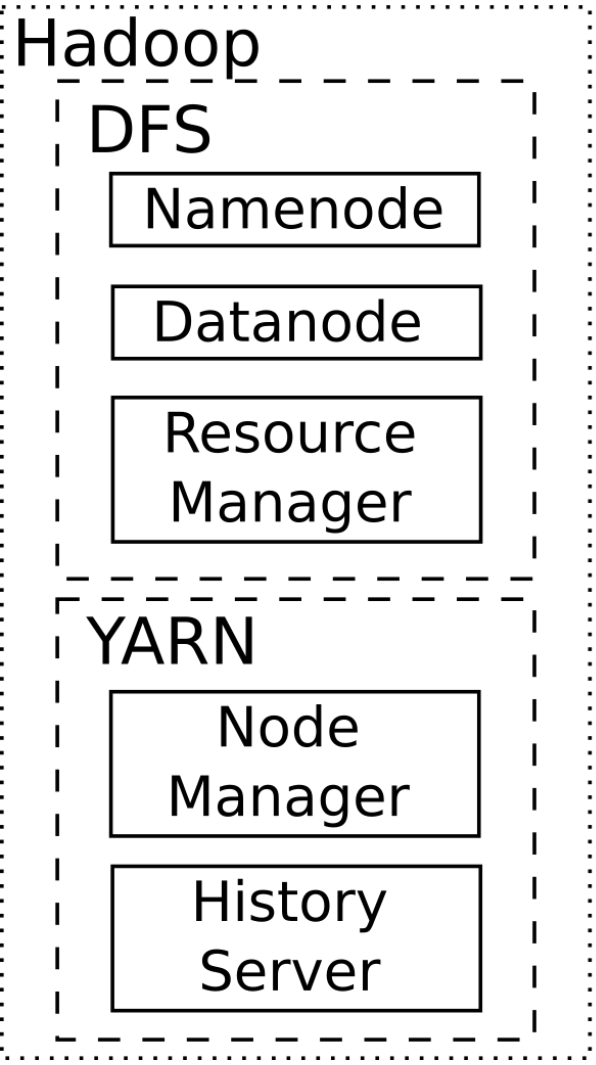

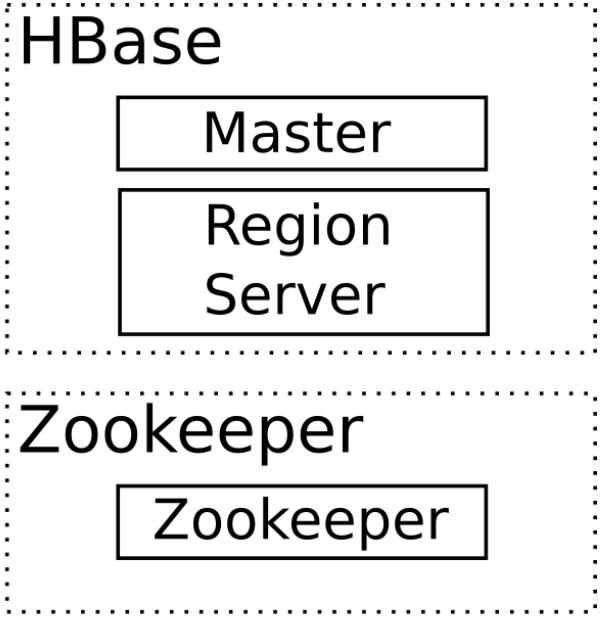

- Hadoop
  - DFS
  - YARN
- HBase
- Zookeeper

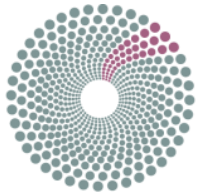

Halyard: BDI Stack (yaml)


```yaml
namenode:
  image: bde2020/hadoop-namenode:1.2.0-hadoop2.8-java8
  container_name: namenode
  networks:
    - hbase
  volumes:
    - ./data/hadoop/namenode:/hadoop/dfs/name
  environment:
    - CLUSTER_NAME=test
  ports:
    - "50070:50070"
  env_file:
    - ./hadoop.env
```

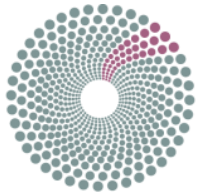

Halyard: BDI Stack: Running

Simply execute the command:

docker-compose up -d

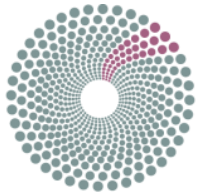

Halyard: BDI Stack: Simple Test


```
$ docker exec -it hbase /bin/bash

$ hbase shell

> list

> create ‘t1’, ‘f1’
```

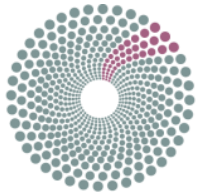

Adding Halyard to BDI Stack


```
FROM bde2020/hadoop-base:1.2.0-hadoop2.8-java8 as hadoop-base
FROM bde2020/hbase-base:1.0.0-hbase1.2.6 as hbase-base
FROM openjdk:8
MAINTAINER Ivan Ermilov <ivan.s.ermilov@gmail.com>
ENV HADOOP_VERSION=2.8.0
COPY --from=hadoop-base /opt/hadoop-$HADOOP_VERSION /opt/hadoop-$HADOOP_VERSION
RUN ln -s /opt/hadoop-$HADOOP_VERSION/etc/hadoop /etc/hadoop
ENV PATH /opt/hadoop-$HADOOP_VERSION/bin:$PATH
ENV HBASE_VERSION=1.2.6
COPY --from=hbase-base /opt/hbase-$HBASE_VERSION /opt/hbase-$HBASE_VERSION
RUN ln -s /opt/hbase-$HBASE_VERSION/conf /etc/hbase
ENV PATH /opt/hbase-$HBASE_VERSION/bin:$PATH
ENV HALYARD_VERSION 1.2
...
```

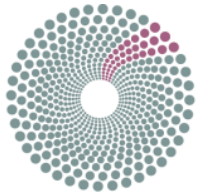

Running Halyard SDK


```
$ docker run -it --rm --network hbase --
 env-file ./hbase.env bde2020/halyard-
 sdk:1.0.0-halyard1.2 /bin/bash
$ ./console
```

Running Halyard SDK


```
> create hbase
> open halyard
> load http://danbri.org/foaf.rdf
```

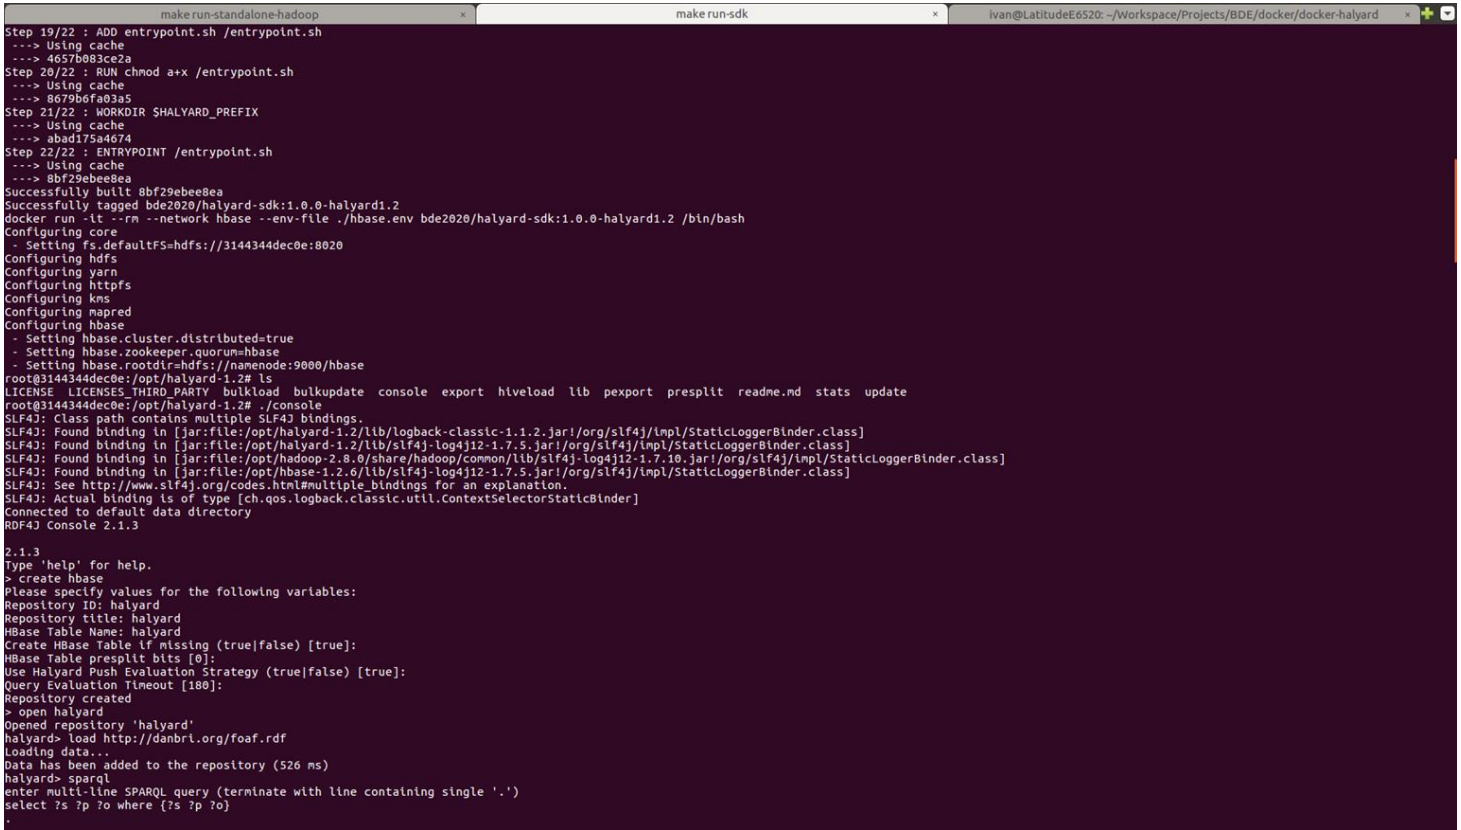

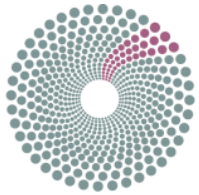

Running Halyard SDK


```
> sparql
select ?s ?p ?o {where ?s ?p ?o} .
.
```

```
make run-standalone-hadoop × | make run-sdk × | ivan@LatitudeE6520: ~/Workspace/Projects/BDE/docker/docker-halyard ×

halyard> sparql
enter multi-line SPARQL query (terminate with line containing single '.')
select ?s ?p ?o where {?s ?p ?o}
.
Evaluating SPARQL query...
+----------------------+----------------------+----------------------+
| s                    | p                    | o                    |
+----------------------+----------------------+----------------------+
| _:node1bltsosg9x41   | rdfs:resource        | "http://perso.hirlimann.net/~ludo/foaf.rdf"@en|
| <http://www.glandscape.com/cgi-bin/prnxml.rss>| rss:title            | "Channel of Filth"@en  |
| <http://www.glandscape.com/cgi-bin/prnxml.rss>| rdfs:seeAlso         | <http://www.glandscape.com/cgi-bin/prnxml.rss>|
| <http://www.glandscape.com/cgi-bin/prnxml.rss>| rdf:type             | rss:channel          |
| <http://mmt.me.uk/foaf.rdf#mischa>| :name                | "Mischa Tuffield"@en   |
| <http://mmt.me.uk/foaf.rdf#mischa>| rdf:type             | :Person              |
| _:node1bltsosg9x44   | :name                | "Margaret Hart"@en     |
| _:node1bltsosg9x44   | :img                 | <http://purl.org/net/danbri/2000/06/mh.jpg>|
| _:node1bltsosg9x44   | :mbox                | <mailto:mags@apocalypse.org>|
| _:node1bltsosg9x44   | rdfs:seeAlso         | <http://www.apocalypse.org/~mags/webwho.xrdf>|
| _:node1bltsosg9x44   | rdf:type             | :Person              |
| _:node1bltsosg9x20   | :name                | "Dean Jackson"@en      |
| _:node1bltsosg9x20   | :mbox                | <mailto:dean@w3.org>   |
| _:node1bltsosg9x20   | :mbox                | <mailto:dino@grorg.org>|
| _:node1bltsosg9x20   | :homepage            | <http://www.grorg.org/dean/>|
| _:node1bltsosg9x20   | rdfs:seeAlso         | <http://www.grorg.org/dean/foaf.rdf>|
| _:node1bltsosg9x20   | rdf:type             | :Person              |
| _:node1bltsosg9x20   | :mbox_sha1sum        | "6de4ff27ef927b9ba21ccc88257e41a2d7e7d293"@en|
| _:node1bltsosg9x17   | :name                | "Dan Brickley"@en      |
| _:node1bltsosg9x17   | :img                 | <http://www.geocities.com/danbfan/Ford3.jpg>|
| _:node1bltsosg9x17   | :isPrimaryTopicOf    | <http://www.geocities.com/danbfan/>|
| _:node1bltsosg9x17   | rdf:type             | :Person              |
| _:node1bltsosg9x45   | :name                | "Libby Miller"@en      |
| _:node1bltsosg9x45   | :mbox                | <mailto:libby.miller@bristol.ac.uk>|
| _:node1bltsosg9x45   | :workplaceHomepage   | <http://ilrt.org/>     |
| _:node1bltsosg9x45   | rdfs:seeAlso         | <http://www.libbymiller.com/webwho.xrdf>|
| _:node1bltsosg9x45   | rdf:type             | :Person              |
| <http://website.lineone.net/~steve_c-t/Scientology/Pickets/10-03-2001/damien.jpg>| dc:description       | "An anti-scientology protest"@en|
| <http://website.lineone.net/~steve_c-t/Scientology/Pickets/10-03-2001/damien.jpg>| rdf:type             | :Image               |
| <http://website.lineone.net/~steve_c-t/Scientology/Pickets/10-03-2001/damien.jpg>| dc:thumbnail         | <http://website.lineone.net/~steve_c-t/Scientology/Pickets/10-03-2001/damien.jpg>|
| <mailto:daniel.brickley@bristol.ac.uk>| dct:isReplacedBy     | <mailto:danbri@danbri.org>|
| _:node1bltsosg9x50   | :name                | "Pastor N Pizzor"@en   |
| _:node1bltsosg9x50   | :made                | <http://www.glandscape.com/cgi-bin/prnxml.rss>|
| _:node1bltsosg9x50   | rdf:type             | :Person              |
| _:node1bltsosg9x19   | :mbox                | <mailto:libby@asemantics.com>|
| _:node1bltsosg9x19   | :mbox                | <mailto:libby.miller@bristol.ac.uk>|
| _:node1bltsosg9x19   | rdf:type             | :Person              |
| _:node1bltsosg9x15   | :accountServiceHomepage| <http://www.flickr.com/>|
| _:node1bltsosg9x15   | :accountProfilePage  | <http://www.flickr.com/people/danbri/>|
| _:node1bltsosg9x15   | :accountName         | "danbri"@en            |
| _:node1bltsosg9x15   | rdf:type             | :OnlineAccount         |
| _:node1bltsosg9x49   | :name                | "Martin L Poulter"@en  |
| _:node1bltsosg9x49   | :mbox                | <mailto:m.l.poulter@bristol.ac.uk>|
| _:node1bltsosg9x49   | :knows               | _:node1bltsosg9x50     |
| _:node1bltsosg9x49   | rdf:type             | :Person              |
| _:node1bltsosg9x13   | :accountServiceHomepage| <http://del.icio.us>   |
| _:node1bltsosg9x13   | :accountProfilePage  | <http://del.icio.us/danbri/>|
| _:node1bltsosg9x13   | :accountName         | "danbri"@en            |
| _:node1bltsosg9x13   | rdf:type             | :OnlineAccount         |
| _:node1bltsosg9x33   | :name                | "Jan Grant"@en         |
| _:node1bltsosg9x33   | :depiction           | <http://ioctl.org/jan/test/wizard.jpg>|
| _:node1bltsosg9x33   | :mbox                | <mailto:jan.grant@bristol.ac.uk>|
```

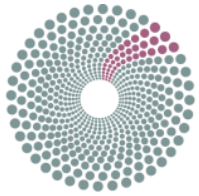

Halyard: BDI Stack (complete)

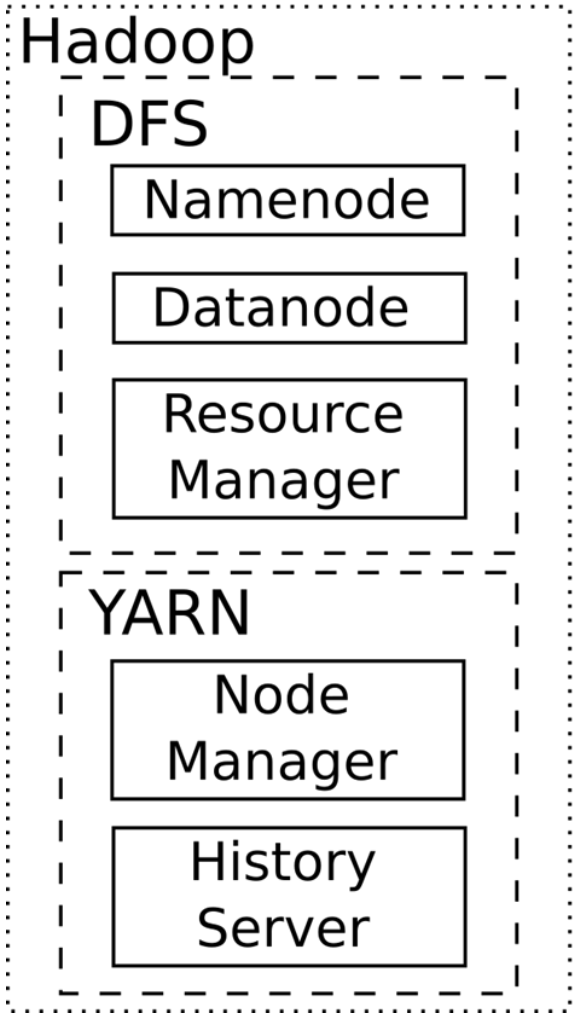

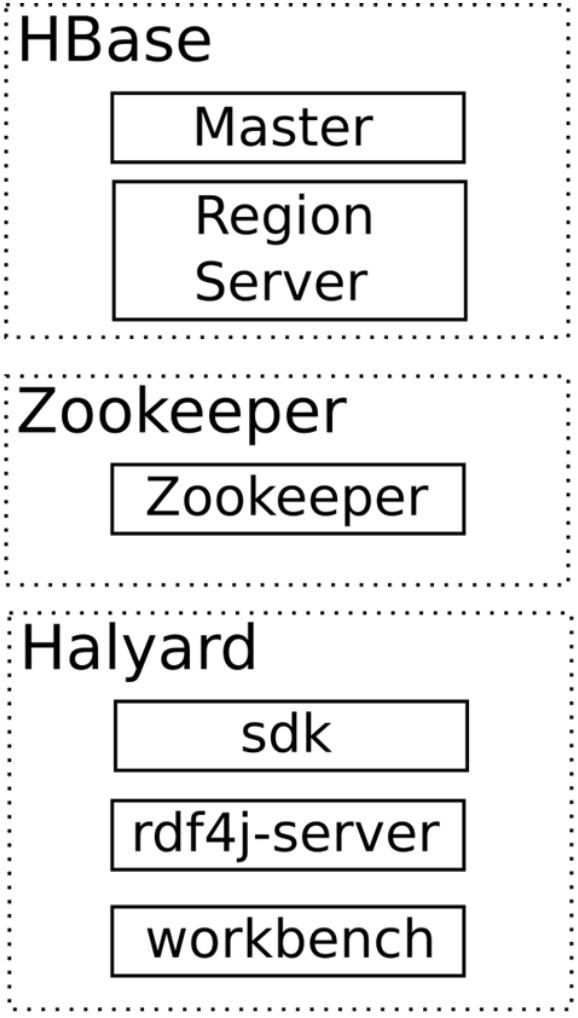

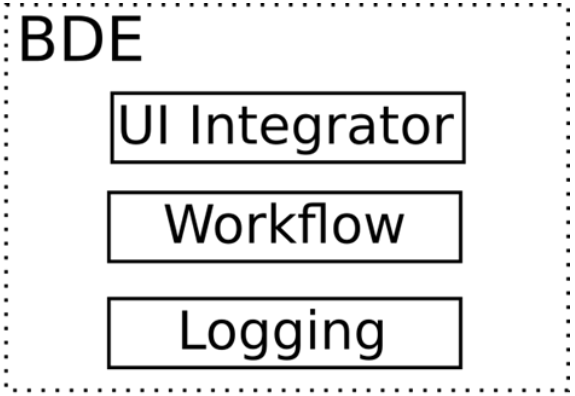

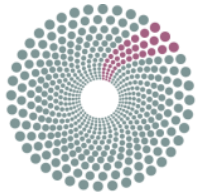

Thank you

Questions?

Github: https://github.com/earthquakesan
@AKSW: http://aksw.org/IvanErmilov.html
Email: iermilov@informatik.uni-leipzig.de
Twitter: @earthquakesan
LinkedIn: https://www.linkedin.com/in/iermilov/